

# MP Out-Of-Order

Presented by Scott & Ian  
Memes: Ethan/Udit/Jason



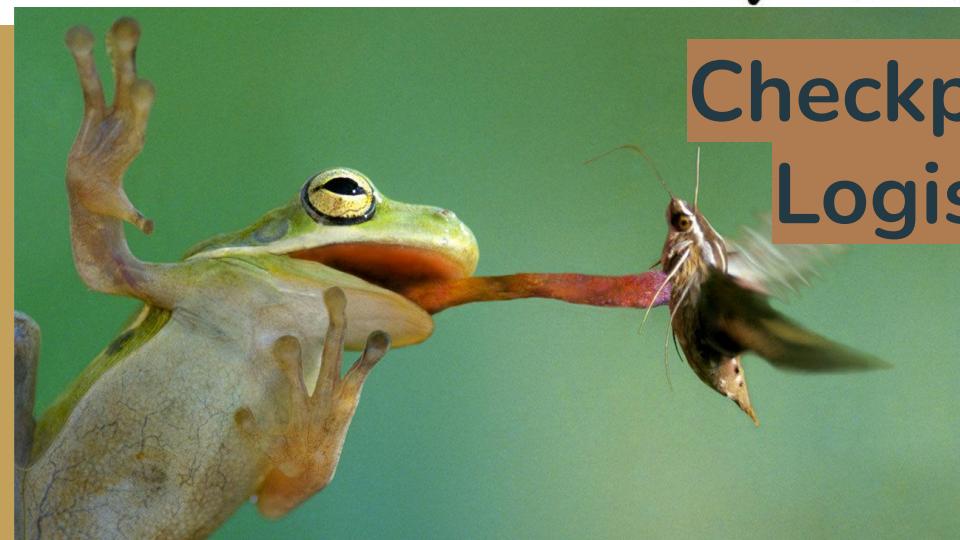
- CP1 deliverables expectations [Scott]
  - Design
    - Interface with memory
    - Each stage
    - Be able to trace every instruction, all arrows need high level field specifications
  - Queue
    - Queue can push & Pop
    - Queue indicates when full
  - Fetch
    - Magic memory, into a queue that stores instruction and PC
- How to make a Queue/Fetch [Ian]
  - How to make head/tail ptr queue, maybe say why it's better than shifting (power)
  - How to integrate it with fetch
  - Expected interface (push/pop)
- Tomasulo Design (briefly) [Scott]
  - Show a basic design (from lecture)
  - Explain verbally what is missing that they would want to add to their diagram
  - Control/Speculative execution
    - Both: wait until control instruction is head of ROB
    - Tomasulo: Flush everything
  - Nomenclature - dispatch vs issue
  - Memory - pretend it's another function unit, with just a single reservation station that doesn't deallocate until commit/retire
    - 2 steps
    - Stores commit in order when they are the head of the ROB
    - Loads get the correct data - youngest store older than the load to the same address
- ERR [Ian]
  - High level block description
  - Go through example instruction execution
  - Memory instructions - complications due to ERR - can't be unified with other instructions
  - Control Flush everything, copy RRF into RAT, reset Free List head/tail pointers.

It is Wednesday.



my dudes

## Checkpoint 1 Logistics



Gentlemen, it is with great pleasure I inform you that



today is wednesday.

# CP1 Deadline Information (3/25 11:59pm)

## Requirements:

- *Design [5]*: Design a high-level block diagram detailing the datapath of your OoO
  - Must be digital - use draw.io or a similar tool. Cannot be handwritten
- *Hardware Submission [12.5]*:
  - Implement a parameterized FIFO with variable depth and width [5]
  - Implement the fetch stage. Must be integrated with magic memory and your parameterized FIFO [7.5]
    - PC initializes to 0x60000000
- Roadmap + Progress Report [2.5]

Total of 20 points.

**Schedule a meeting with your assigned mentor TA to demonstrate the above!!!!**



# How to: FIFO

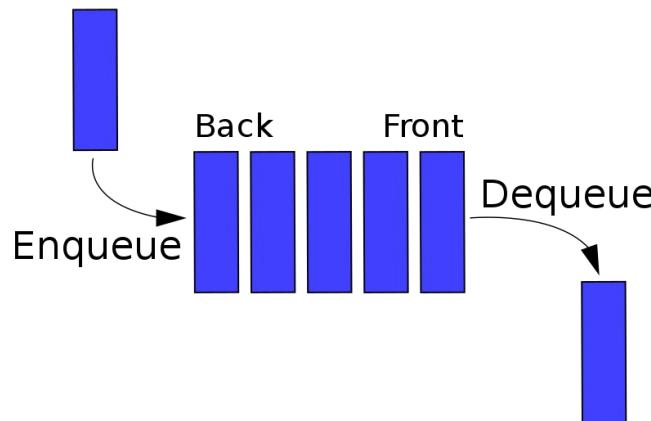
alamy

alamy

alamy

# What is a Queue/FIFO?

First In, First Out

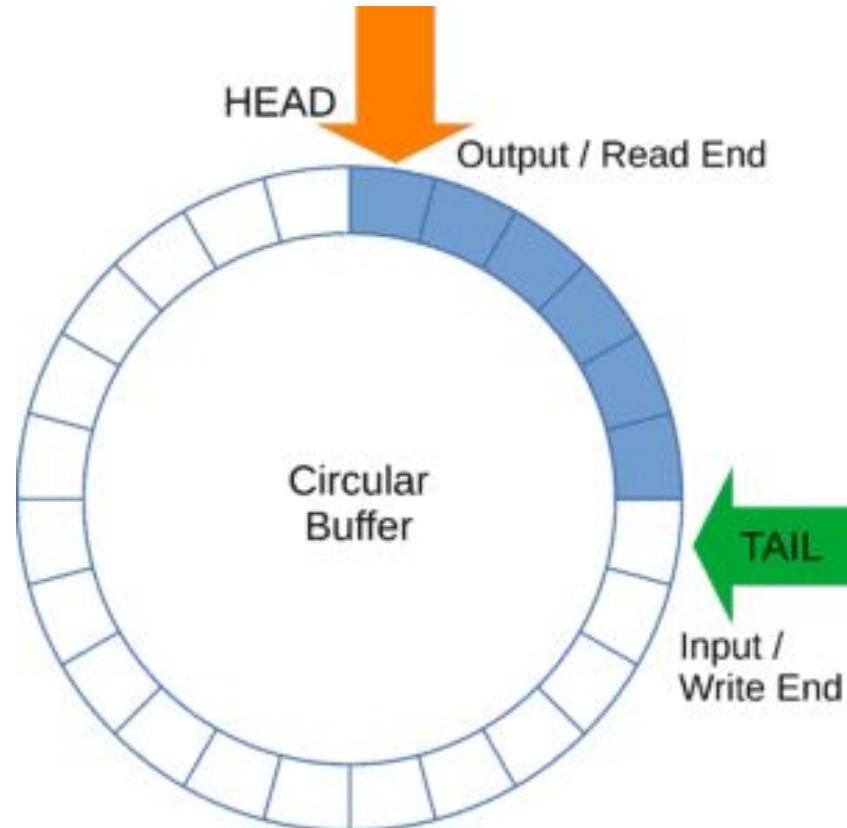


# What does a FIFO look like in hardware?

Big buffer (registers or sram) store data

Internal head and tail pointers

- Read from head, increment head when dequeuing
- Write to tail, increment tail when enqueueing



# How do we control this FIFO?

**input enqueue** - Control signal to enqueue **enqueue\_wdata**

**input enqueue\_wdata** - Input data to enqueue

**input dequeue** - Control signal to dequeue **dequeue\_rdata**

**output dequeue\_rdata** - Output data to dequeue

Are we missing anything?

# When is the FIFO full or empty?

```
logic [DATA_WIDTH-1:0] entries [16];  
  
logic [3:0] head_ptr;  
logic [3:0] tail_ptr;
```

Head\_ptr == Tail\_ptr → We know that the FIFO is either full or empty

```
assign full = (head_ptr == tail_ptr);  
assign empty = (head_ptr == tail_ptr);
```

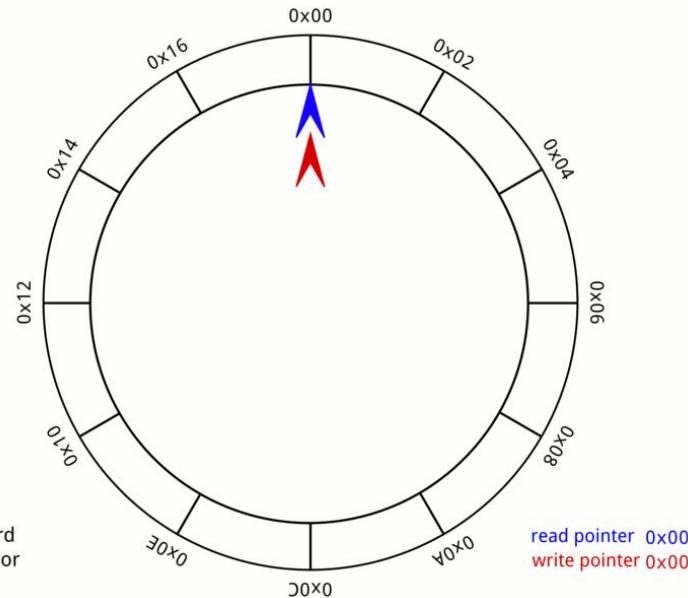
How do we distinguish between the two?

## Full and Empty (cont)

Detect whether the tail “lapped” the head.

The tail\_ptr would *overflow* before reaching the head pointer

Let's add another bit!



Example for a circular buffer found on wikipedia. Head is read pointer, tail is write pointer.

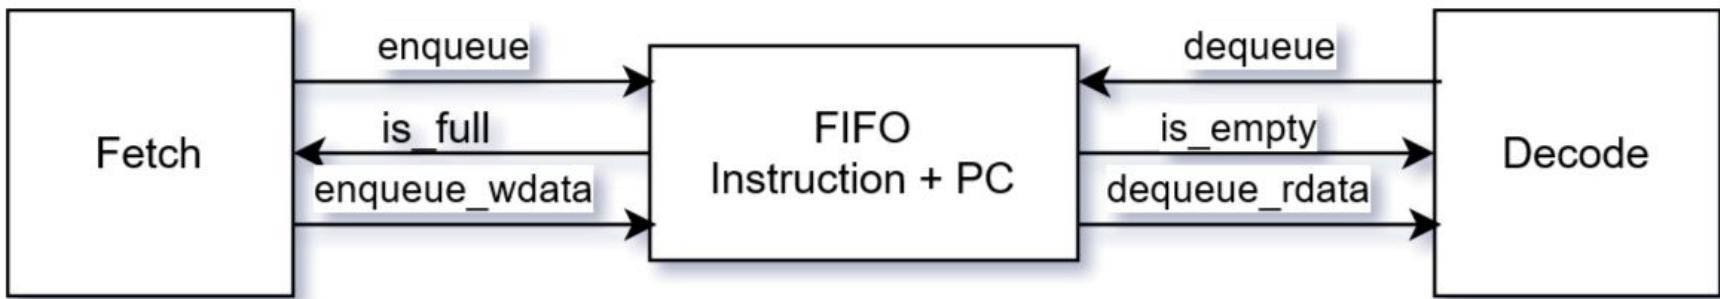
## Full and Empty (cont cont)

```
logic [DATA_WIDTH-1:0] entries [16];  
  
logic [4:0] head_ptr;  
logic [4:0] tail_ptr;
```

With that additional bit,  
how do we detect full?

```
assign full = (  
    . . . (head_ptr[3:0] == tail_ptr[3:0]) &&  
    . . . (head_ptr[4] != tail_ptr[4])  
)
```

# Fetch to Decode



Make sure fetch stalls when the FIFO is full!  
Why did we do this?

A close-up photograph of a green tree frog, likely a species of Phyllomedusa, perched on a light-colored, textured surface. The frog's body is mostly bright green with darker, mottled patterns on its back and legs. Its large, dark eyes are prominent. The background is blurred.

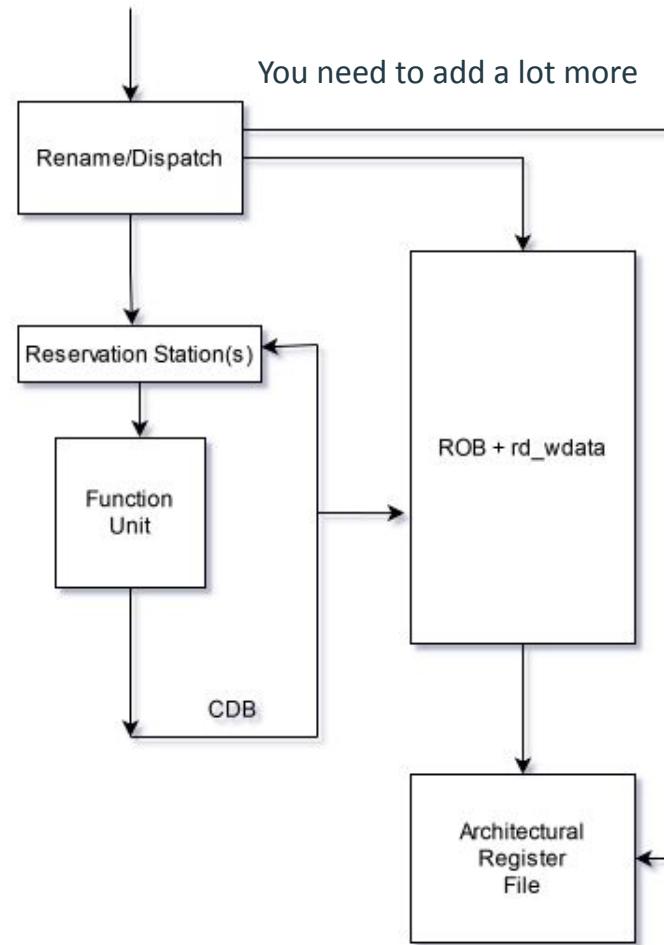
# Tomasulo's Algorithm

# Tomasulo's Algorithm & OOO



## Motivations:

- Tomasulo leverages instruction-level parallelism to achieve higher performance than in-order pipelines
  - Tomasulo's algorithm can schedule instructions without dependencies in parallel or otherwise out-of-order.
- Tomasulo is relatively simpler (conceptually) compared to Explicit Register Renaming (R10k style)
- Simpler designs can be developed faster, leaving more time for debugging and/or advanced features
  - More on the pros of ERR (and how it works) in later slides



# Some Nomenclature

**Fetch:** Grab instructions from memory

**Decode:** Parse the instructions and determine operands, immediates, etc.

**Dispatch:** Put instructions into ROB and RS. Wait for source registers to be ready.

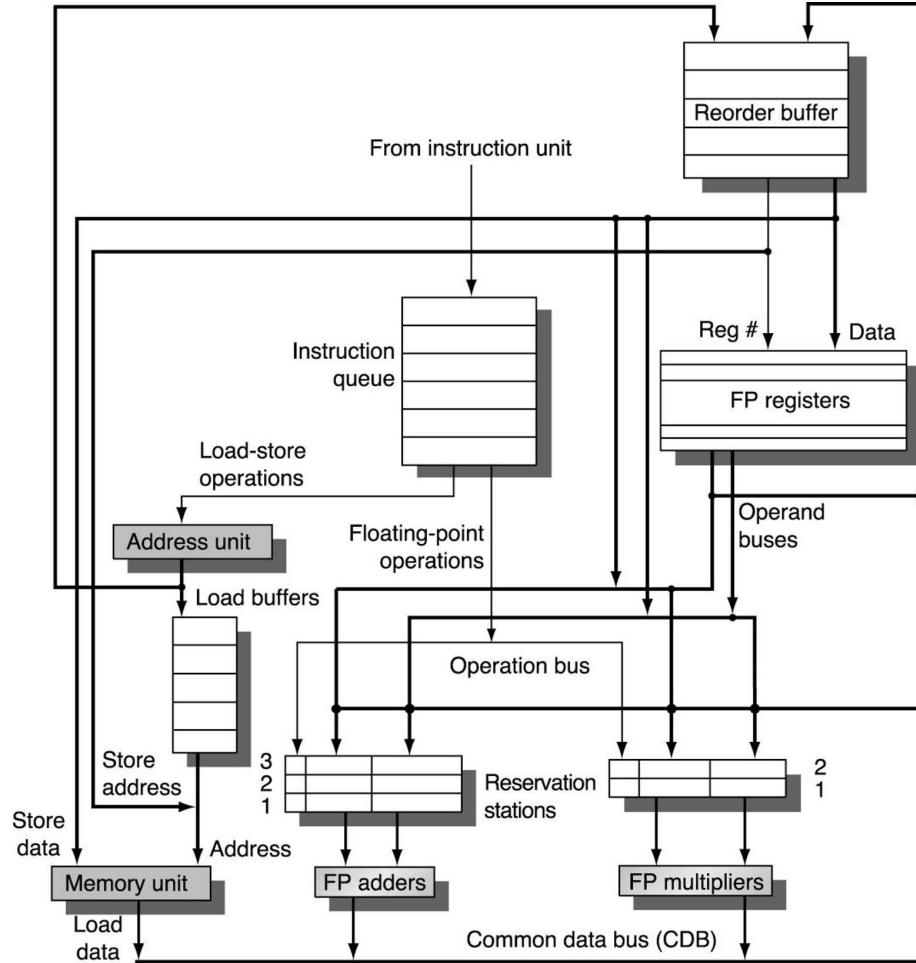
**Issue:** Once the source register values are ready, send the instruction to the appropriate execution unit

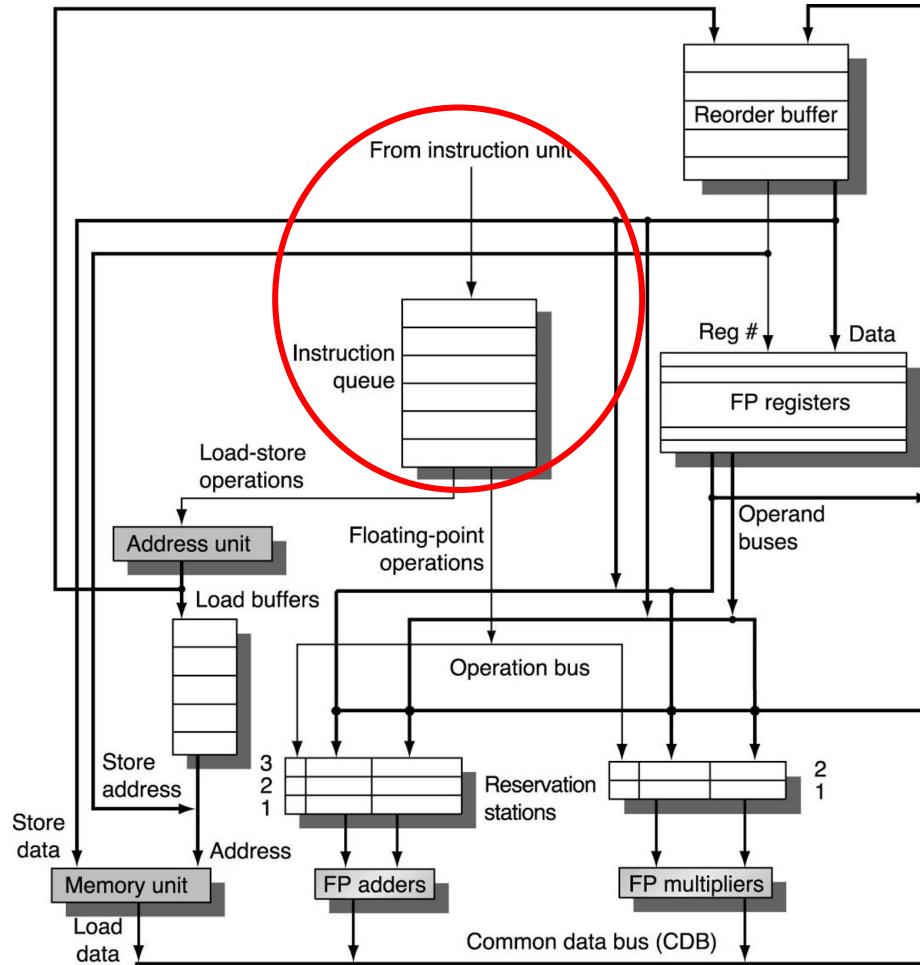
**Execute:** Do the execution (mults, adds, etc.)

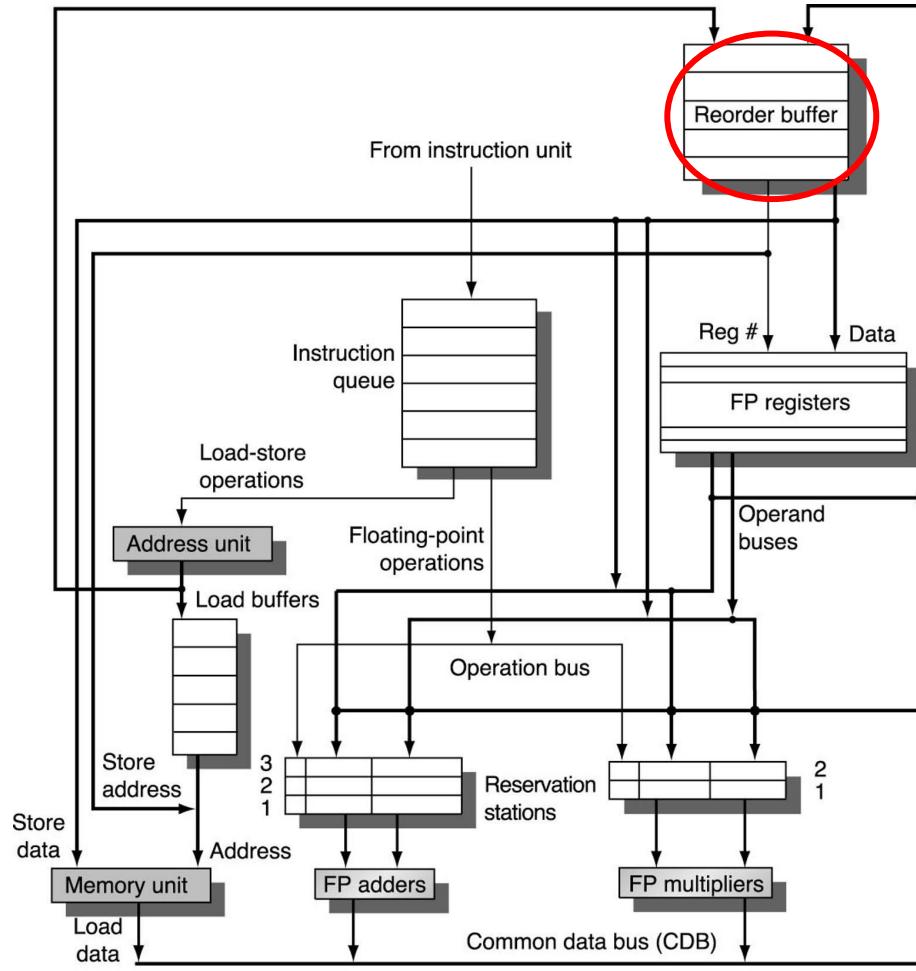
**Writeback:** Broadcast the value calculated during execute on the CDB

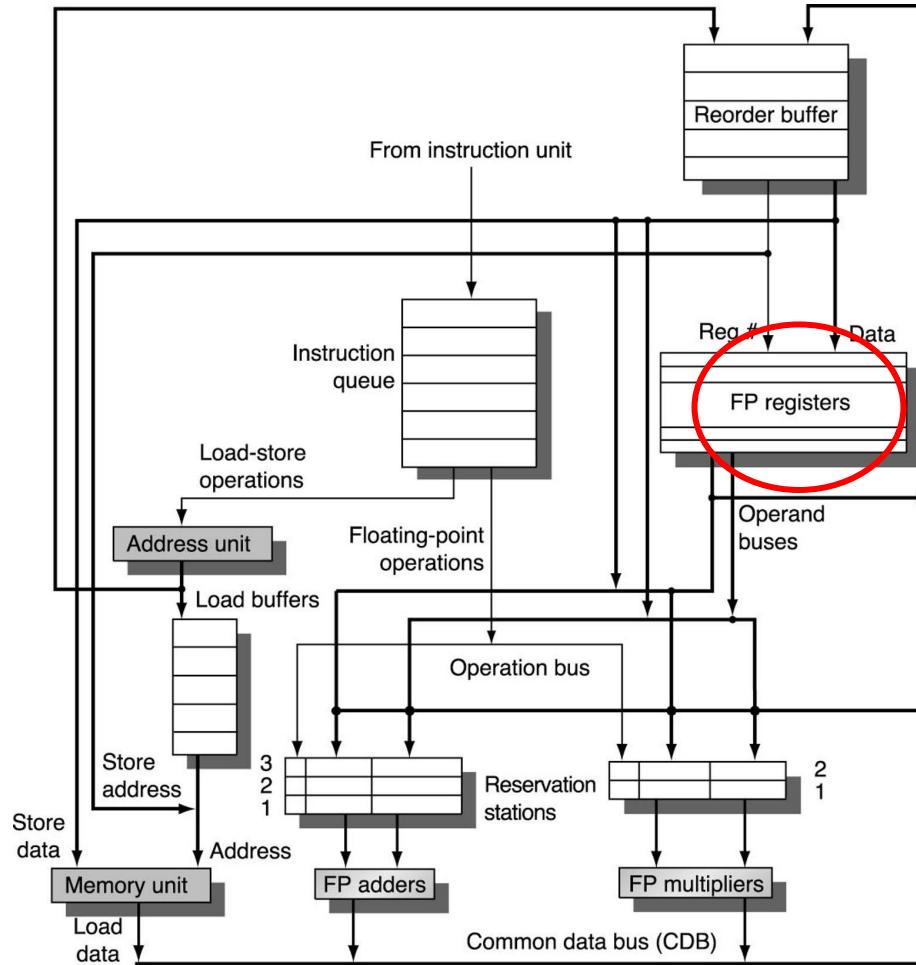
**Commit/Retire:** Write the oldest value in the ROB into the architectural register file (Tomasulo) or update the RRF (explicit renaming). Also store to memory if committing a store instruction.

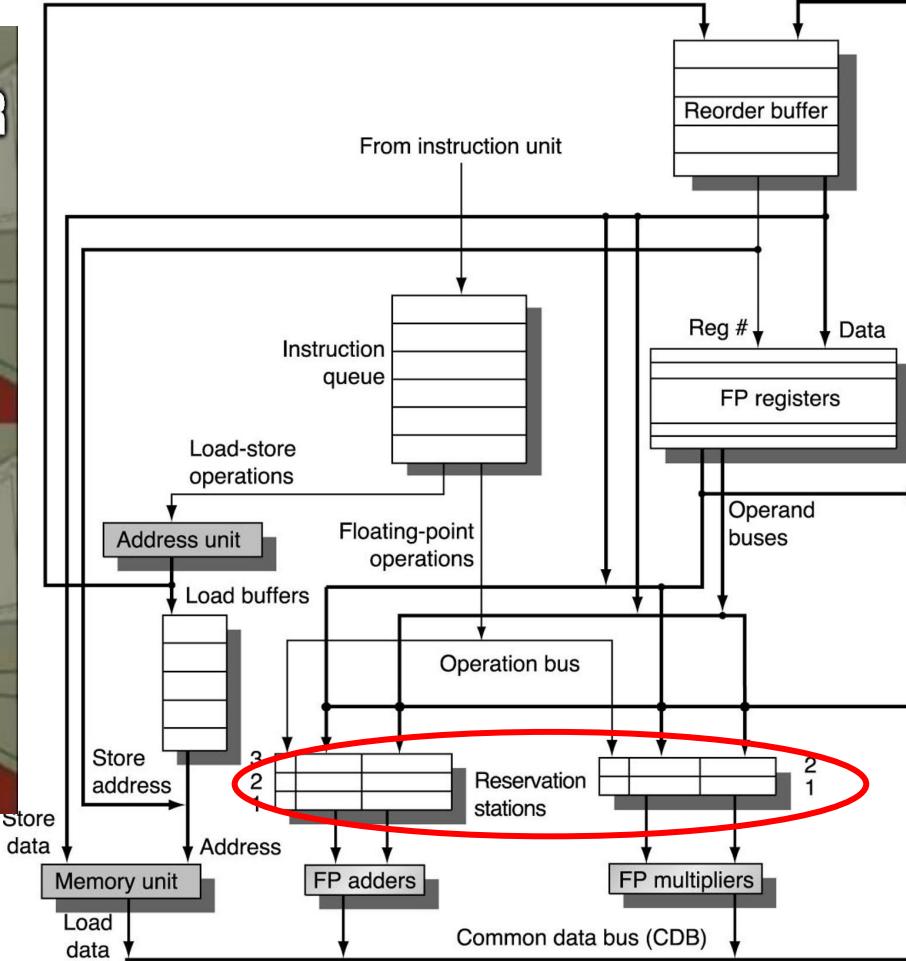
You need to add a lot more

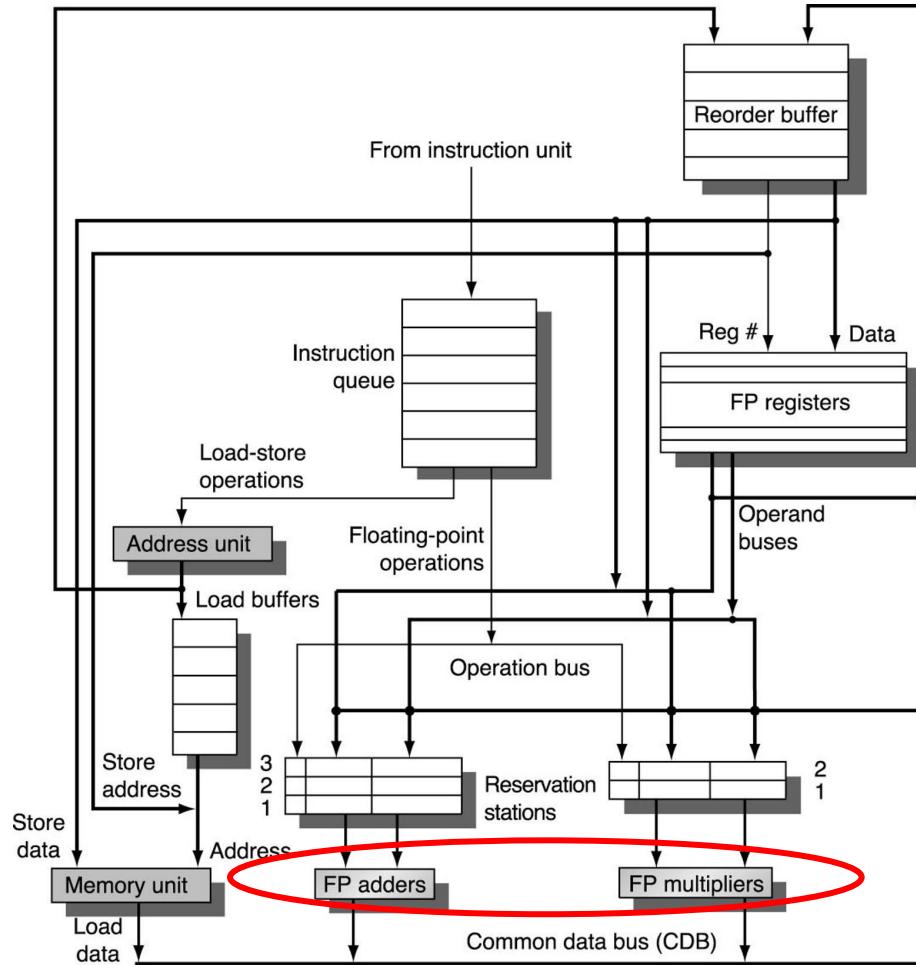


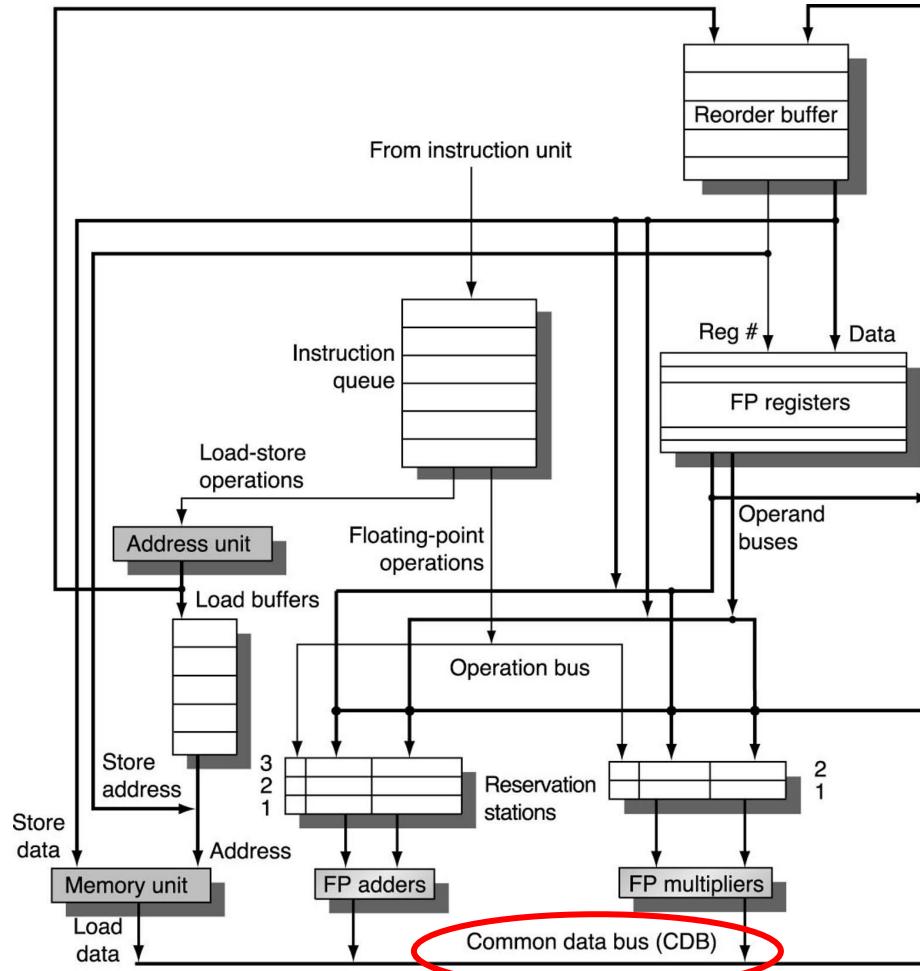


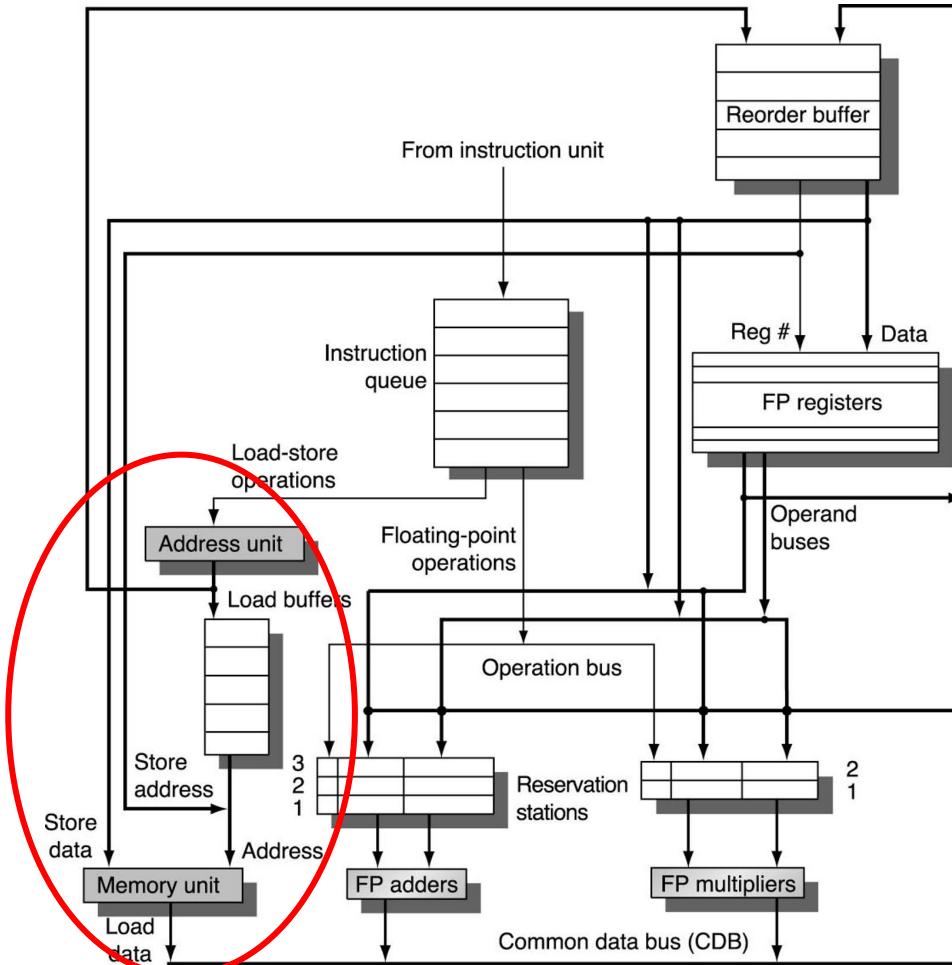




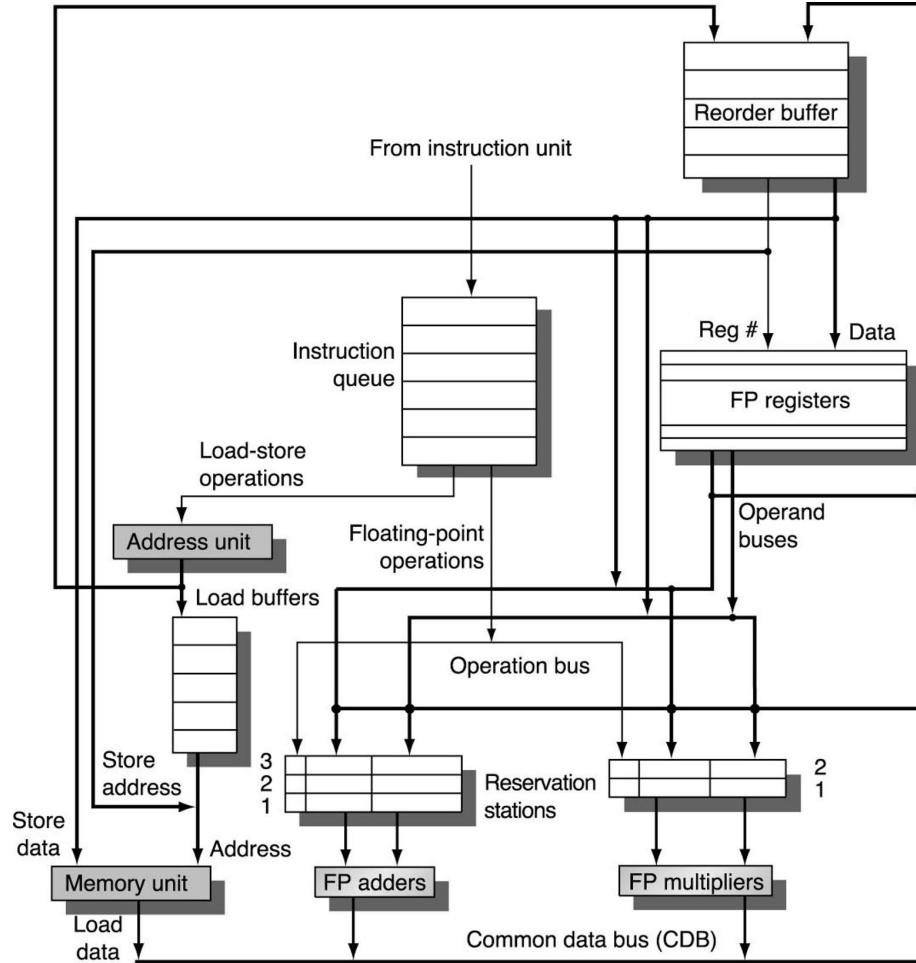






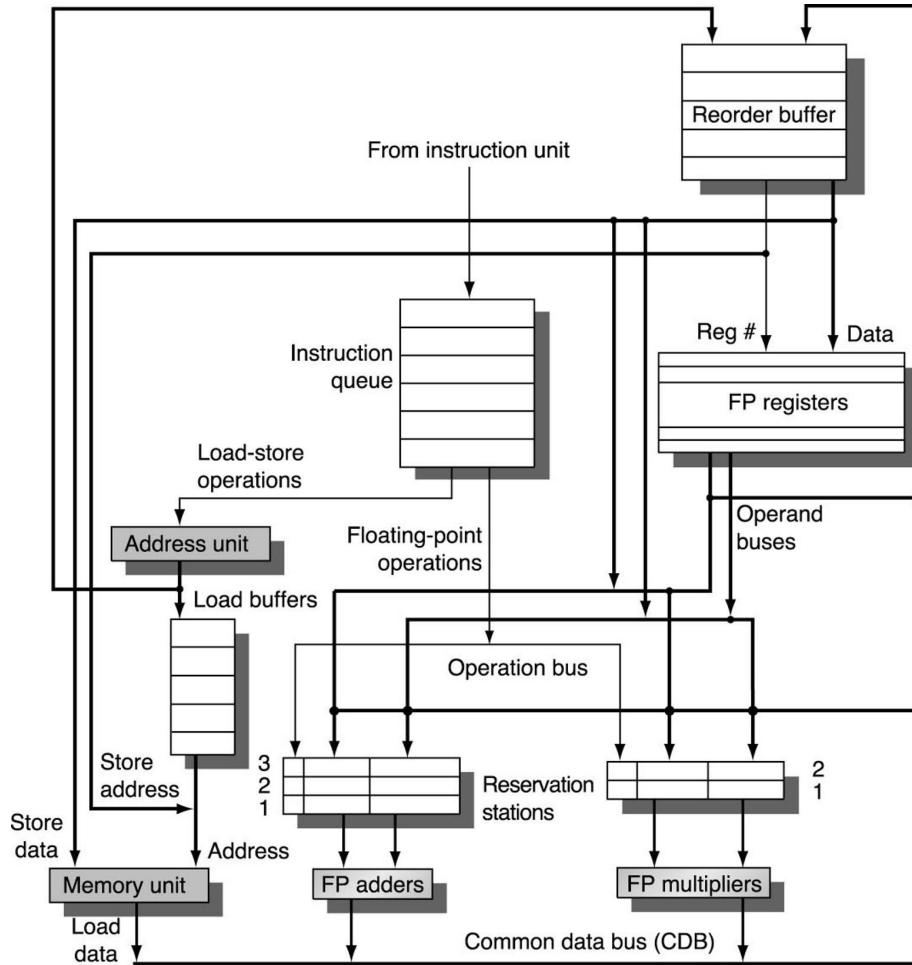


What's missing?



## What's missing?

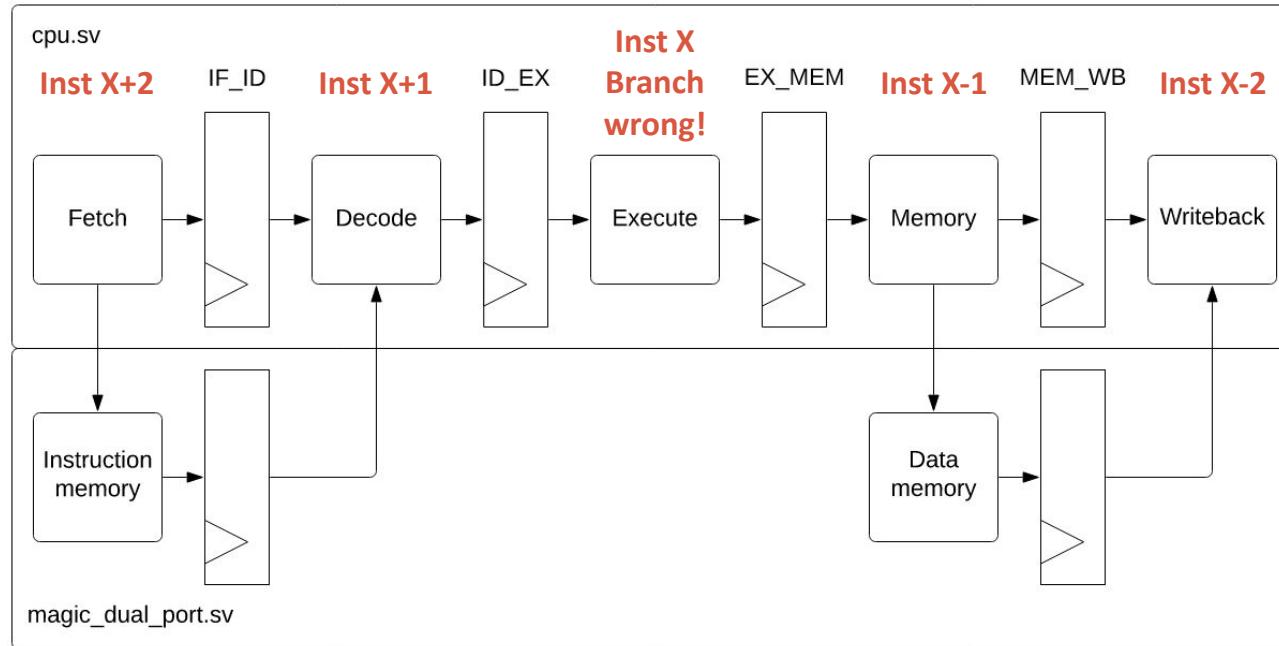
- I cache
- D cache
- Branch prediction and recovery
- (& you are using integer not floating point so reg file and execution units will be different)
- Load/Store logic is flexible and you probably will want something different.



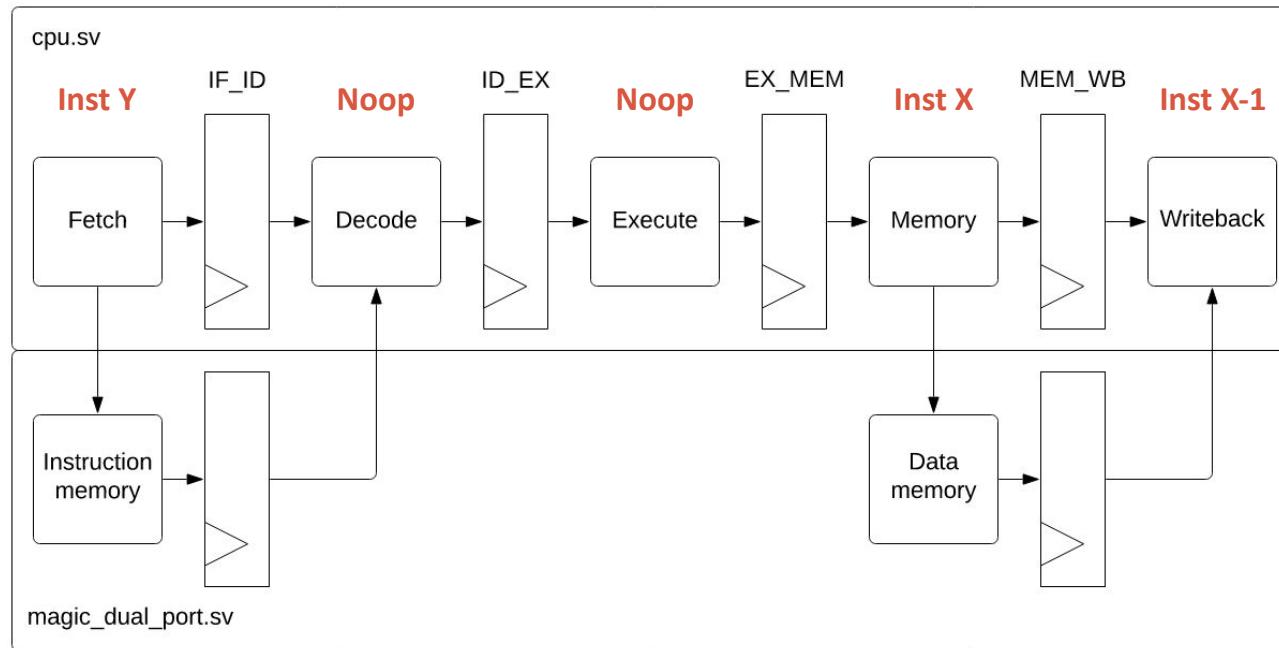


## Branch Misprediction (Speculative Execution)

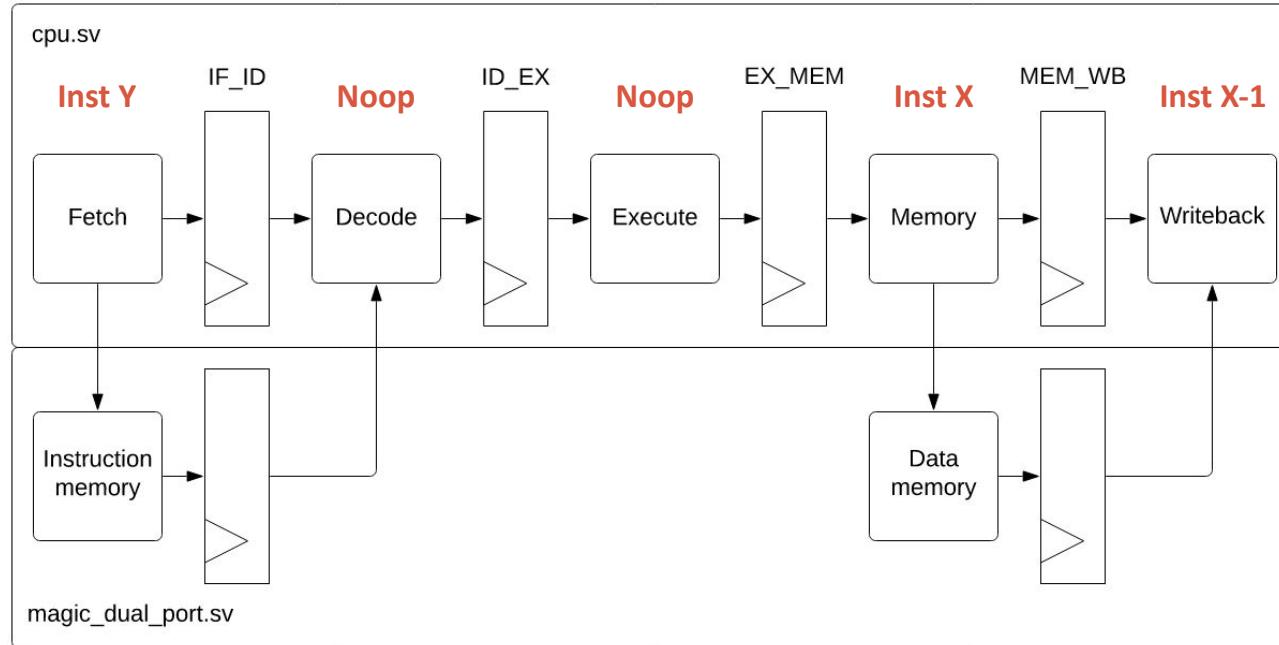
# Speculative Execution: Pipelined Design



# Speculative Execution: Pipelined Design



# Speculative Execution: Pipelined Design



Q: But how do we do this in an out of order processor?

Hint: Which structure maintains program order?

# Speculative Execution: Pipelined Design

**Q:** But how do we do this in an out of order processor?

**V1:** On commit, check if the branch was mispredicted. If it was, squash everything younger than the branch (every other instruction).

[Minimum implementation]

Requires flush signals sent to all of your hardware blocks. These are triggered by the ROB commit logic. This is a little trickier in Explicit Register Renaming (see later slide)



# Speculative Execution: Pipelined Design

**Q:** But how do we do this in an out of order processor?

**V2:** Squash branches immediately once you figure out they were mispredicted.  
[Advanced Feature: Early Branch Recovery]

Requires branch dependency logic to be tagged on every instruction  
Also requires checkpointing processor state before every branch so that it can be recovered on a mispredict.

For more details, speak to your mentor TA.





Memory

# Memory Operations: Correctness

## Rules for memory correctness:

- The contents of memory (including cache) must never be speculative.
- The contents of memory (including cache) must be updated in order
- Loads to address X must reflect the result of the youngest store that modifies address X that is older than the load.

# Memory Operations: Challenges

Handling memory instructions is non-trivial in an out-of-order processor

**Q:** You tell me - why are memory instructions challenging?

# Memory Operations: Challenges

Handling memory instructions is non-trivial in an out-of-order processor

**Q:** You tell me - why are memory instructions challenging?

**A:** Memory disambiguation. Memory operations have register dependencies *and* memory address dependencies! Furthermore, you can't rollback a write to memory once it occurs!

# Memory Operations: Example

## Program:

SW: x10, 12(x15) # Store x10 ->[12 + x15]

ADD: x5, x0, x2

LW: x1, 12(x15) # Load [12 + x15] -> x1

# Memory Operations: Example

## Program:

SW: x10, 12(x15) # Store x10 ->[12 + x15]

ADD: x5, x0, x2

LW: x1, 12(x15) # Load [12 + x15] -> x1



# Memory Operations: Example

## Program:

SW: x10, 12(x15) # Store x10 ->[12 + x15]

ADD: x5, x0, x2

LW: x1, 12(x14) # Load [12 + x14] -> x1

# Memory Operations: Example

## Program:

SW: x10, 12(x15) # Store x10 ->[12 + x15]

ADD: x5, x0, x2

LW: x1, 12(x14) # Load [12 + x14] -> x1



# Memory Operations: Example

## Program:

SW: x10, 12(x15) # Store x10 ->[12 + x15]

ADD: x5, x0, x2

LW: x1, 12(x14) # Load [12 + x14] -> x1



**Need memory disambiguation!**

**Loads/Stores Require Two Steps:**

1. Determine source/destination address
2. Read/Write memory

# Memory Operations: Details

**When Can a Store Execute (Write to memory?)**

**When Can a Load Execute (Read from memory)?**

# Memory Operations: Details

**When Can a Store Execute (Write to memory?)**

A: On Commit

**When Can a Load Execute (Read from memory)?**

V1: Once all older stores are committed

[Minimum implementation]

Easiest implementation is to stall dispatch

- After a store is dispatched, stall dispatching all subsequent memory instructions until that store commits.

# Memory Operations: Details

**When Can a Store Execute (Write to memory?)**

A: On Commit

**When Can a Load Execute (Read from memory)?**

V2: Once all older store addresses are known and conflicting stores are committed

[Advanced Feature]

Need to keep track of relative ordering of loads & stores. Requires complex load-store queue(s)

# Memory Operations: Details

**When Can a Store Execute (Write to memory?)**

A: On Commit

**When Can a Load Execute (Read from memory)?**

V3: Once all older store addresses are known, with store values forwarded to the load before the store commits.

[Bigger advanced feature]

Requires loads to grab values from store queue in addition to load/store queue(s).

# Memory Operations: Details

**When Can a Store Execute (Write to memory?)**

A: On Commit

**When Can a Load Execute (Read from memory)?**

V4: Speculatively decide (before the addresses are known) whether store values should be forwarded to loads and *rollback* if wrong  
[Hardest implementation & worth the most advanced feature points.]

See [Fire & Forget paper](#) or create a custom load-store queue implementation with data dependency predictors and squashing logic.

# Explicit Register Renaming



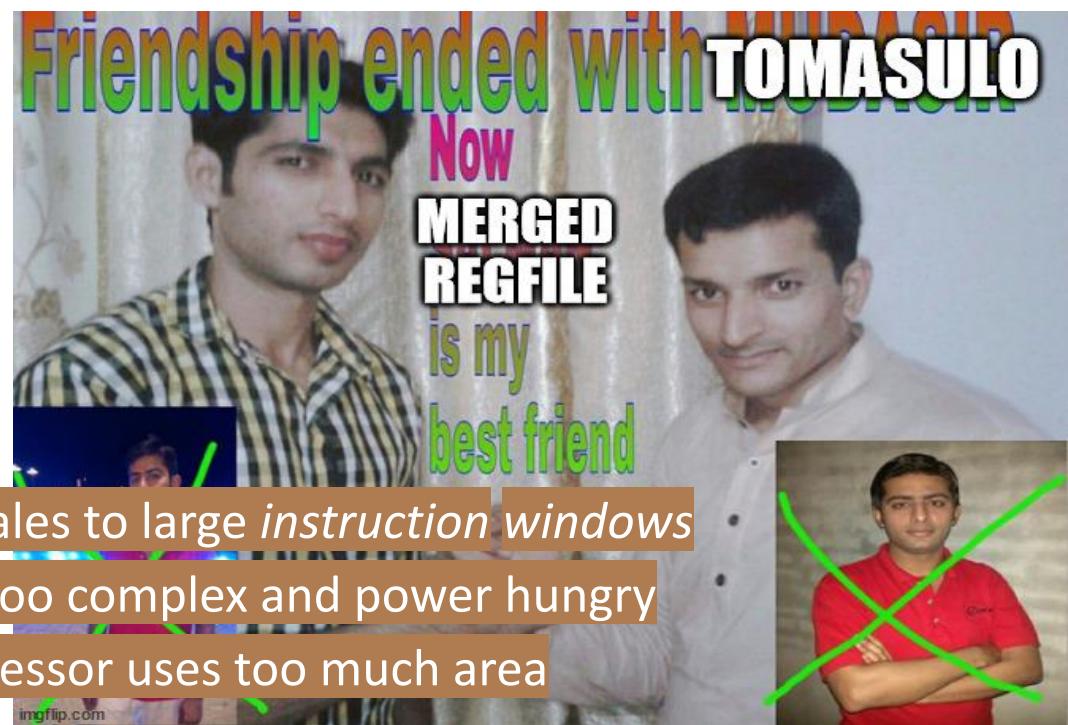
Average Tomasulo fan    Average merged regfile enjoyer

# Explicit Register Renaming (ERR)



Motivations:

- Tomasulo poorly scales to large *instruction windows*
- ROB becomes too complex and power hungry
- The entire processor uses too much area



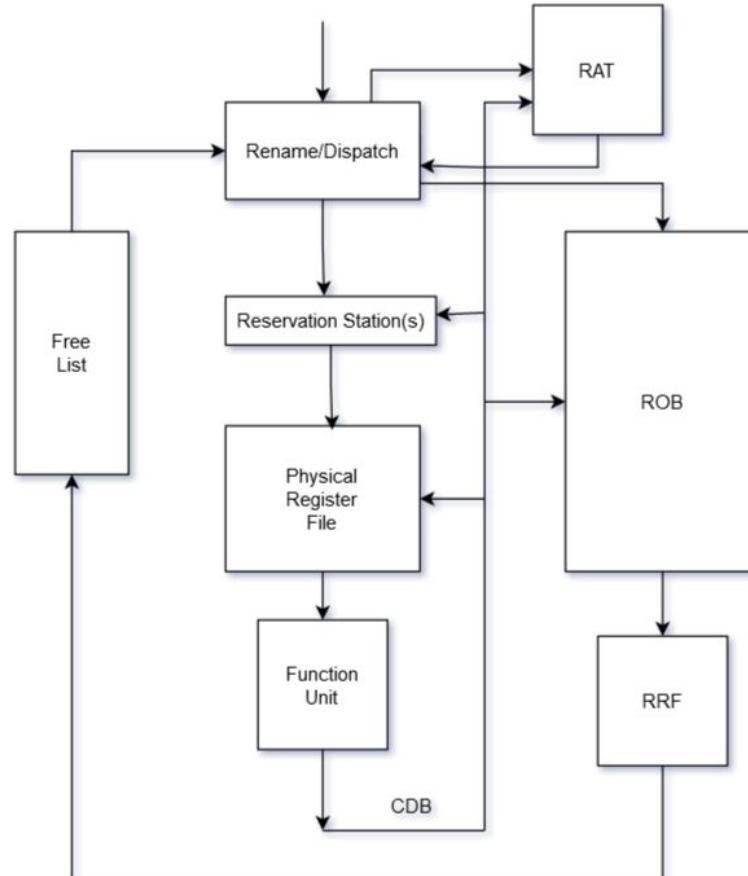
Can we be more efficient?

# Datapath Diagram

Only the Physical Register File stores data

It is only read *after* an instruction leaves a reservation station

If you decide to do ERR, you will need to include much more detail!



# Free List

It's just a FIFO\*!



\*with a bit of fanciness for control instructions

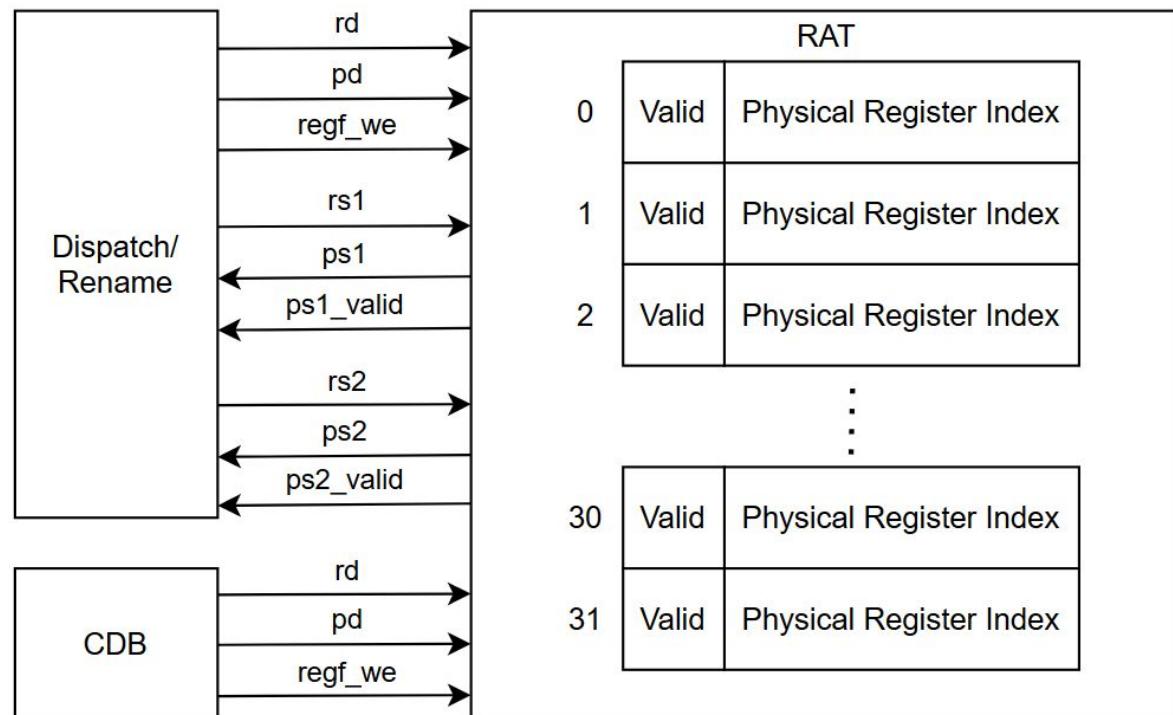
# RAT - Register Alias Table

Dispatch/Rename:

- Maps arch sources ( $rs1/rs2$ ) to phys sources ( $ps1(ps2)$ )
- Renames rd to pd, marking invalid

CDB:

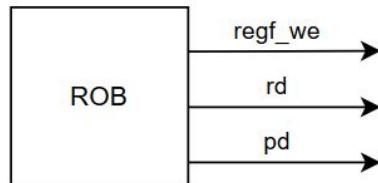
- Sets entry rd to valid if it still maps to pd.



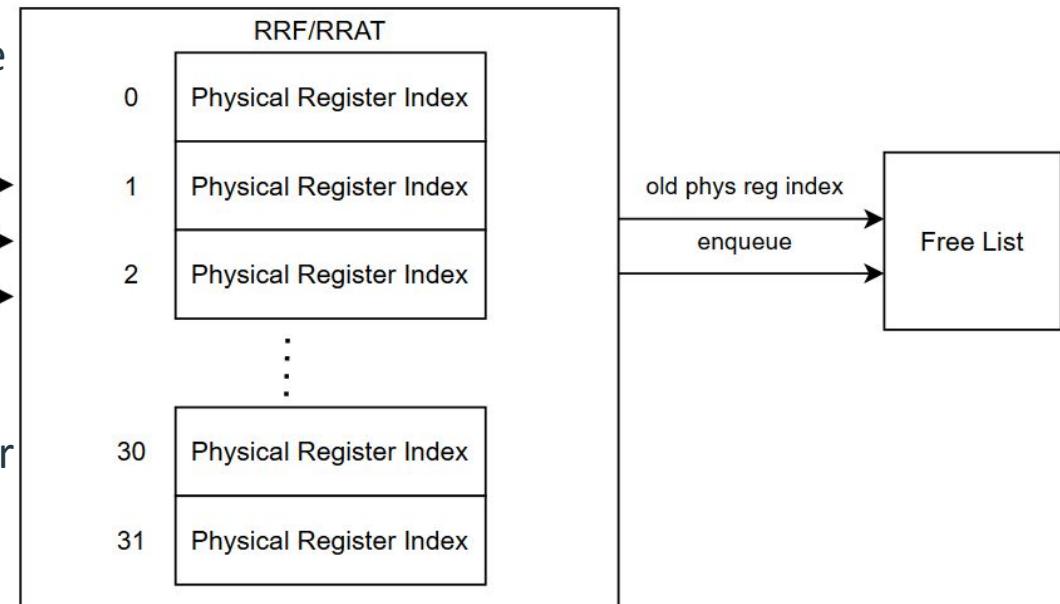
# RRF/RRAT (Retirement RF/RAT)

Similar to RAT - just no valid field

Stores the committed state of the processor



Used to store the architectural state of the processor (needed for speculation)



# Cycle 0:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

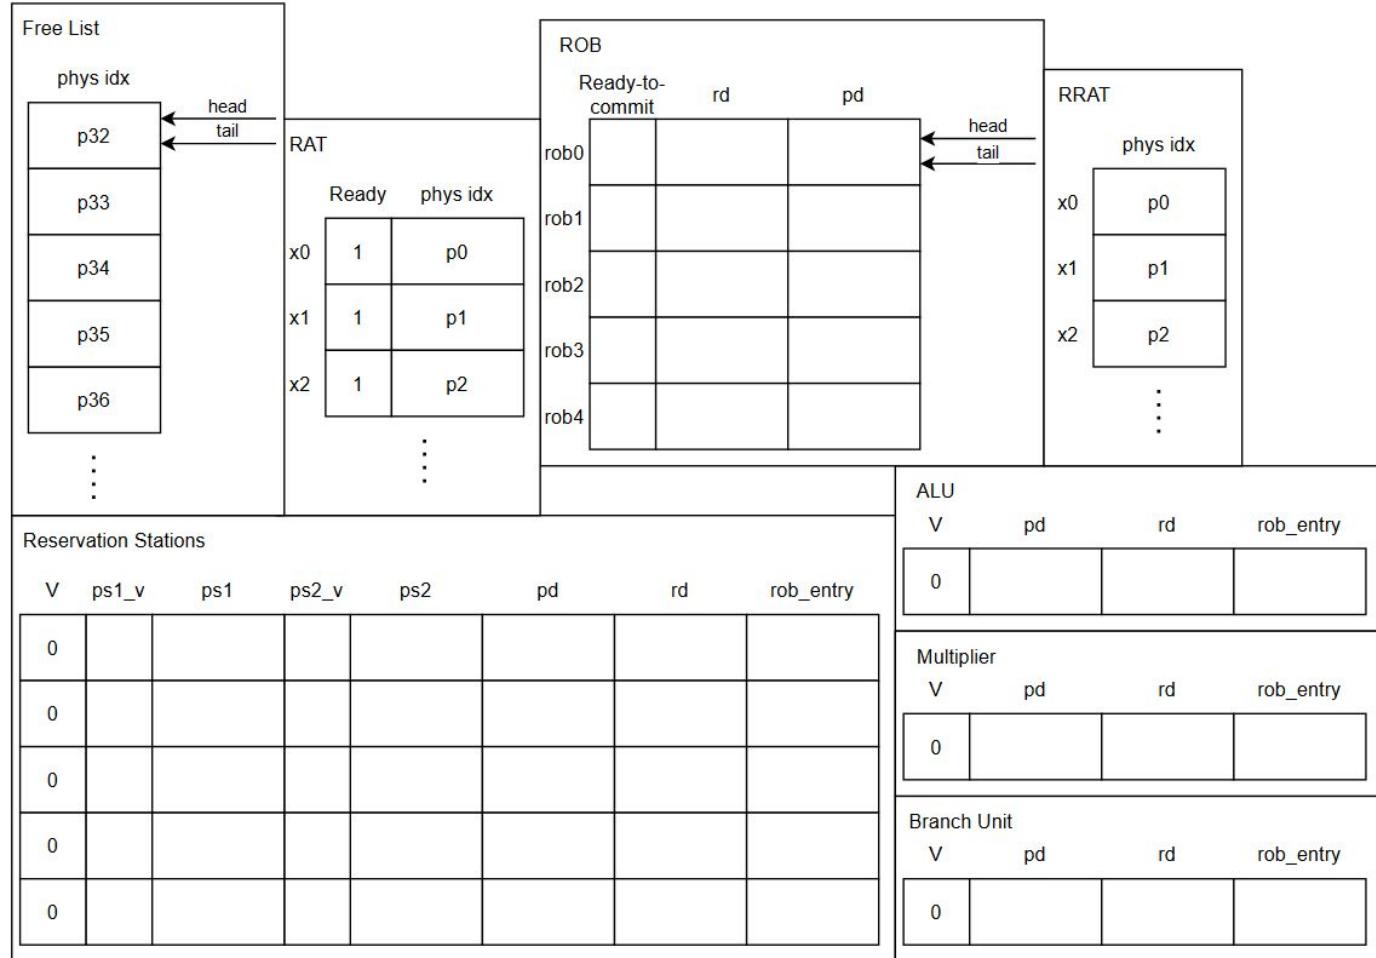
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles



# Cycle 1:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

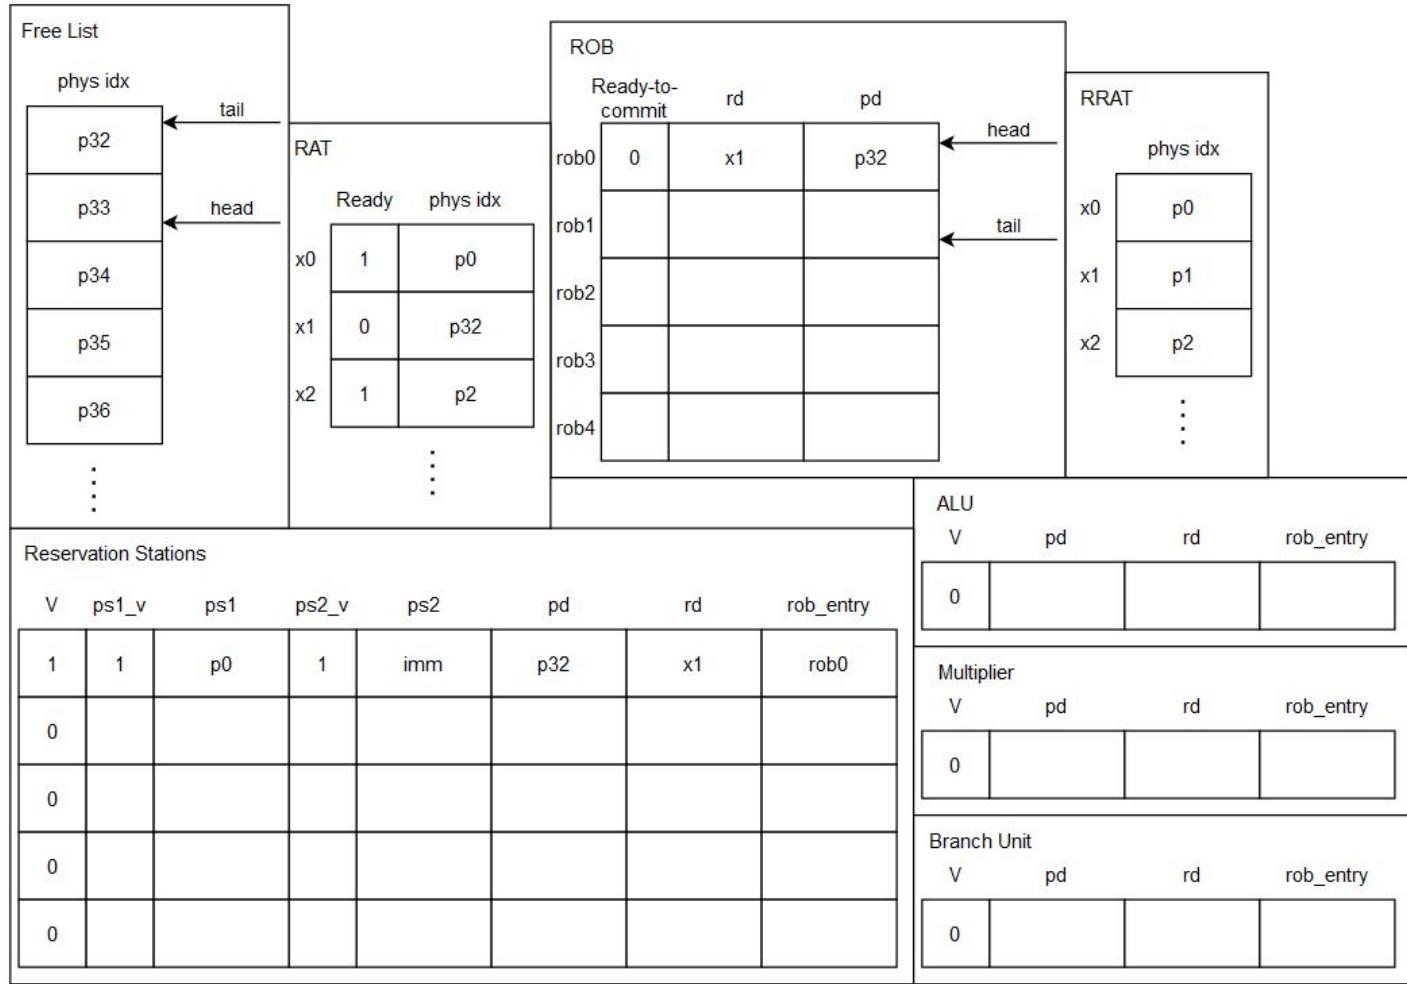
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles



# Cycle 2:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

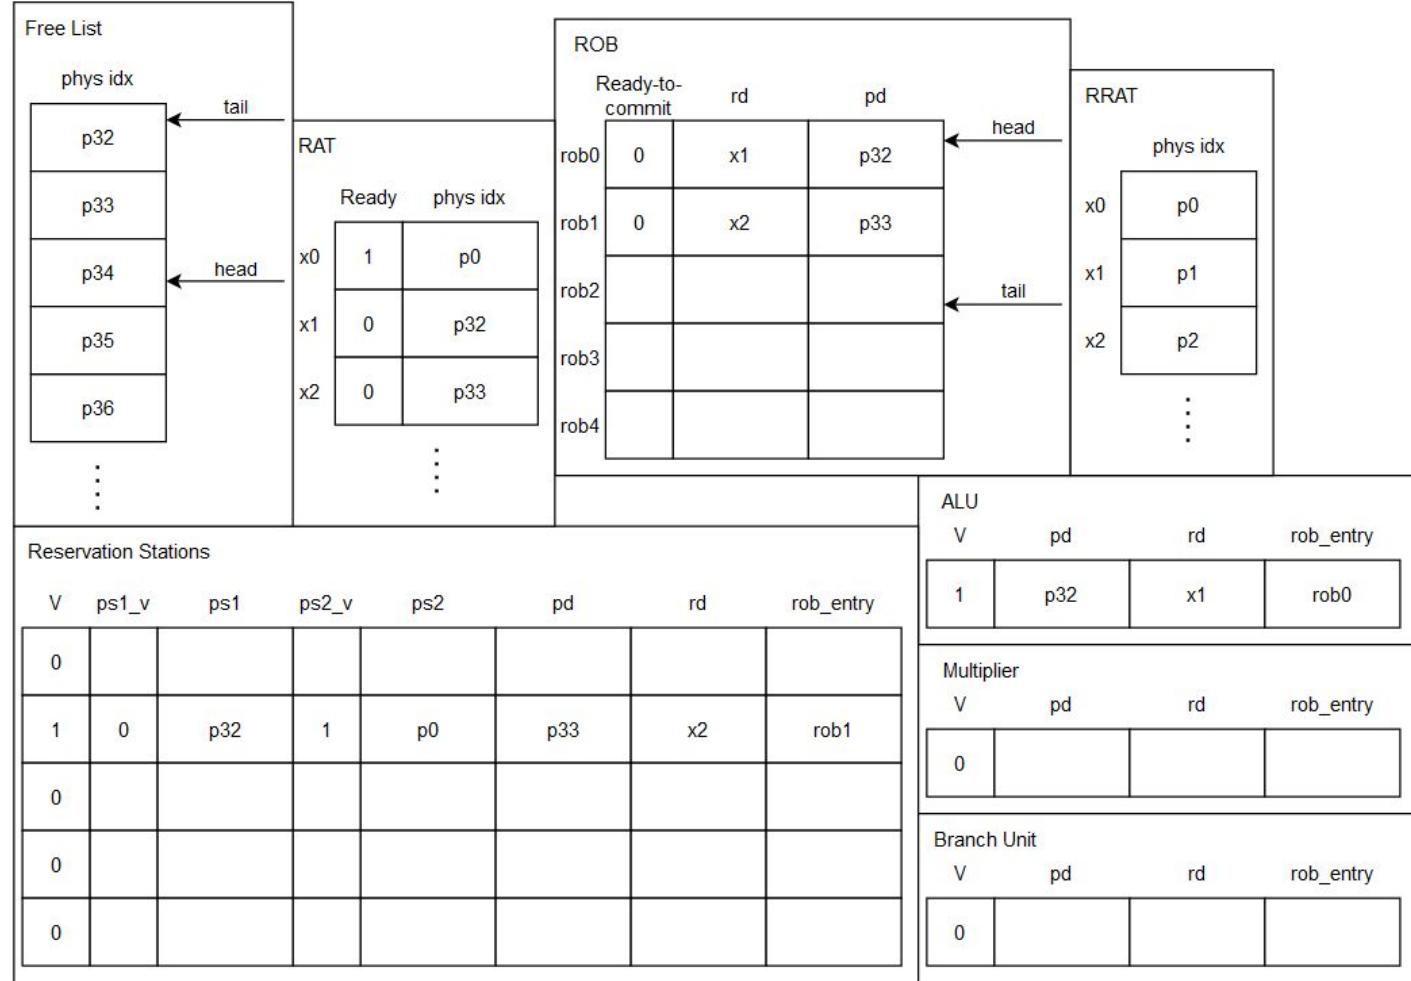
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles



# Cycle 3:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

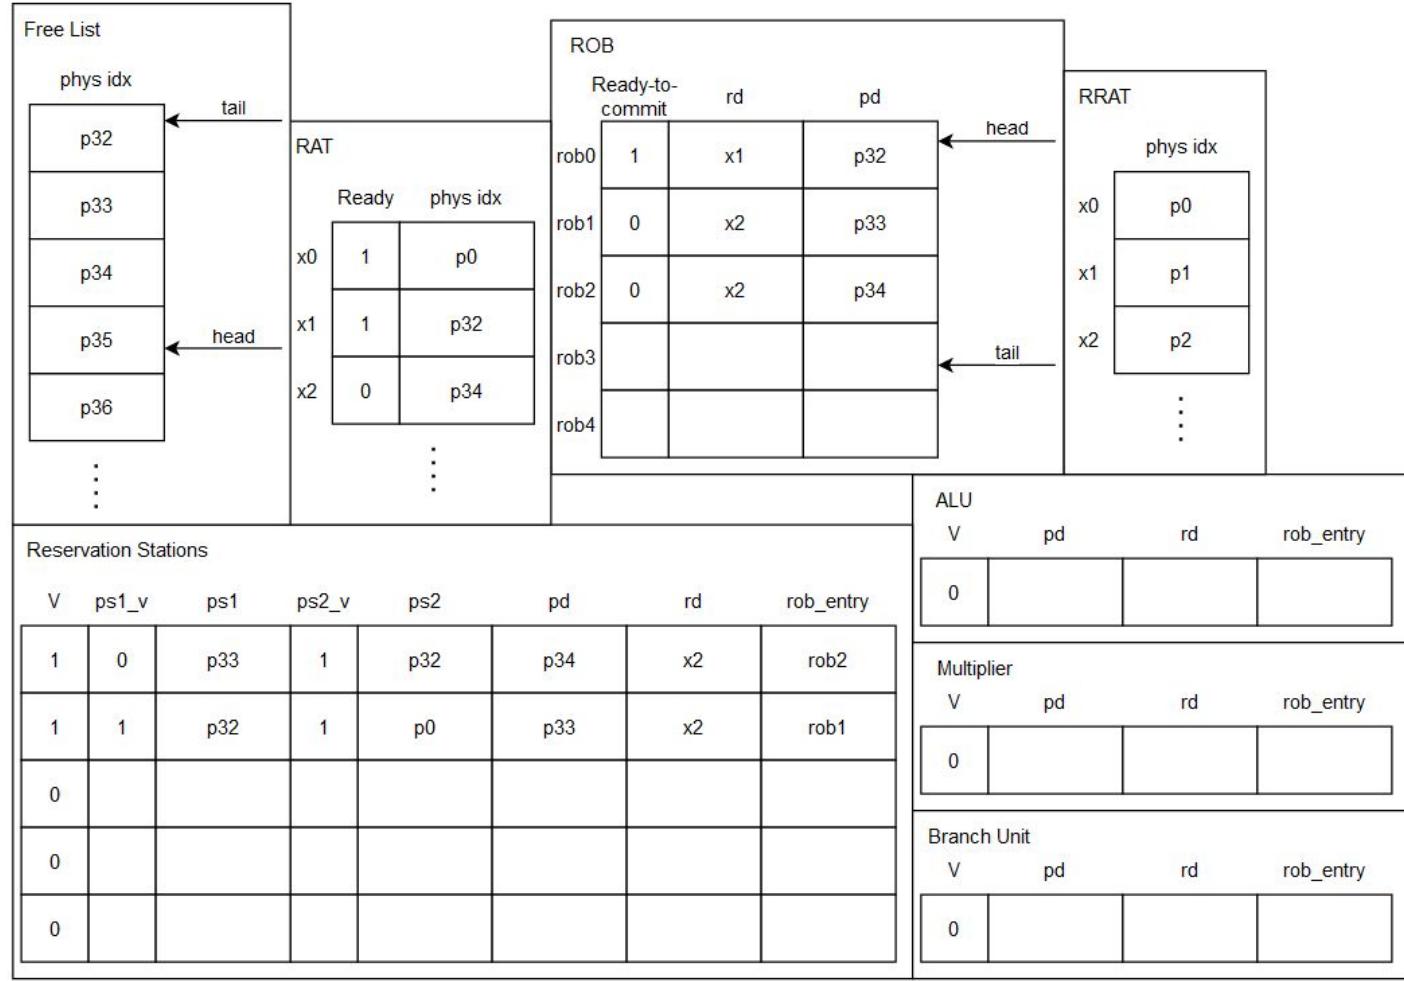
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles



# Cycle 4:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

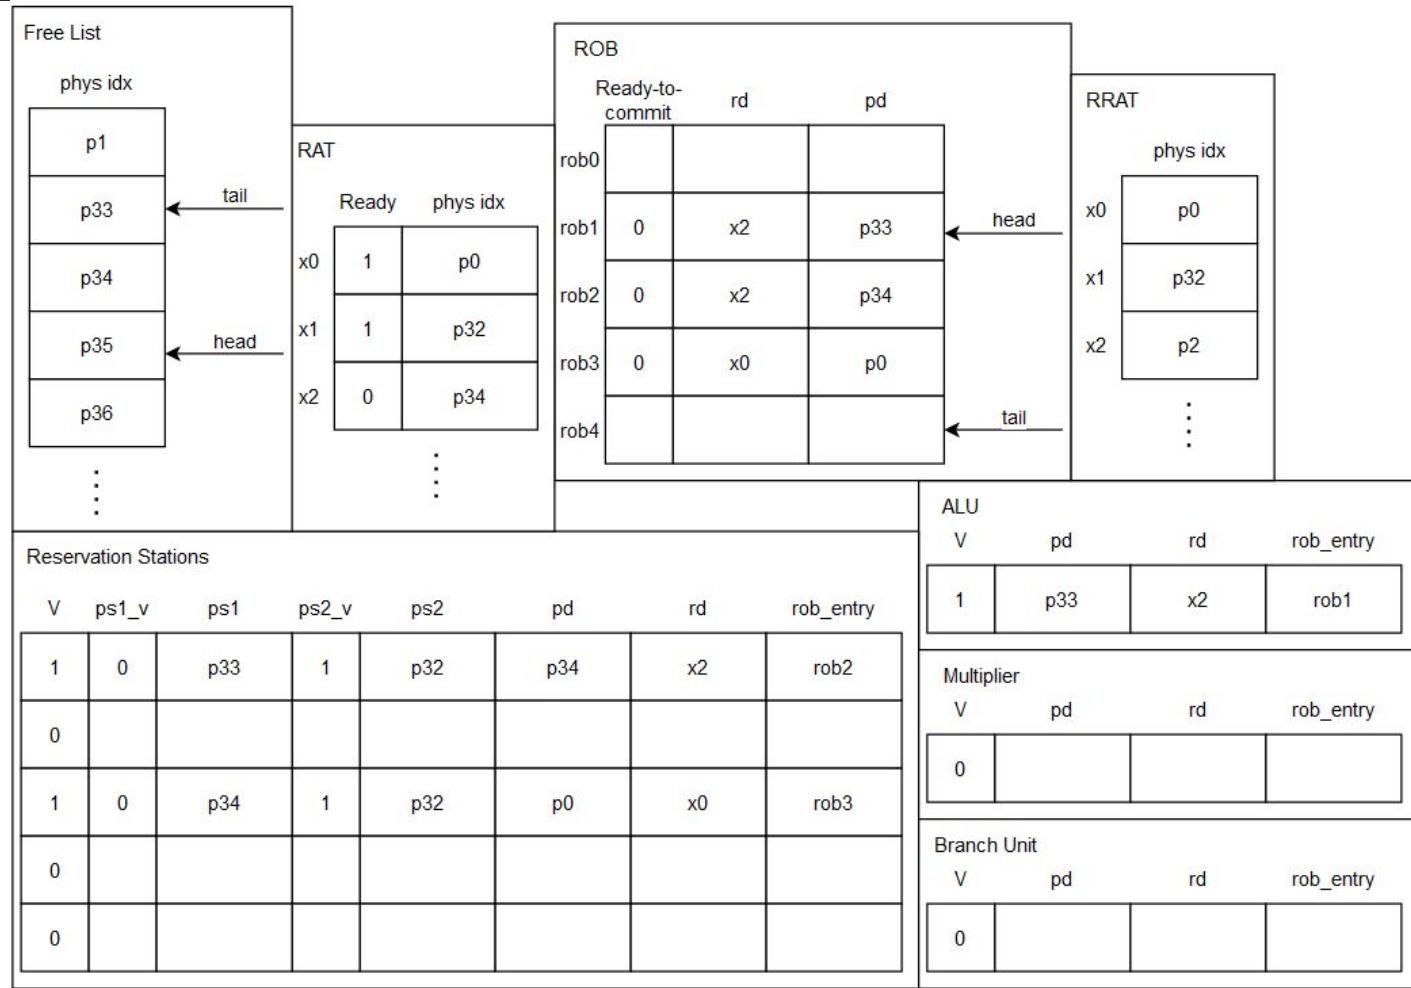
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles



# Cycle 5:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

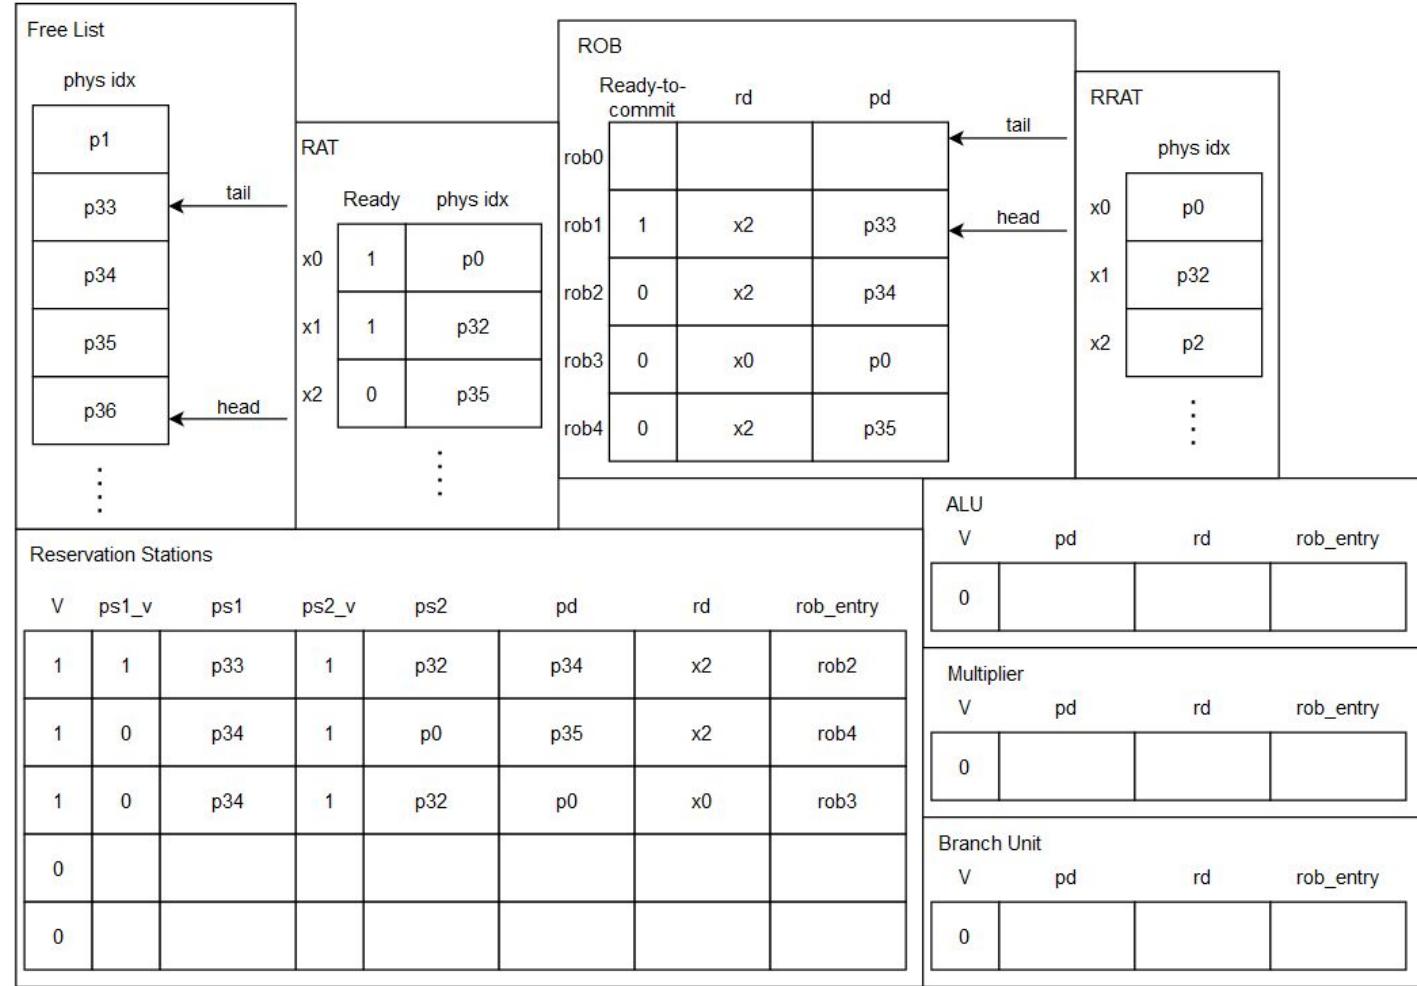
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles

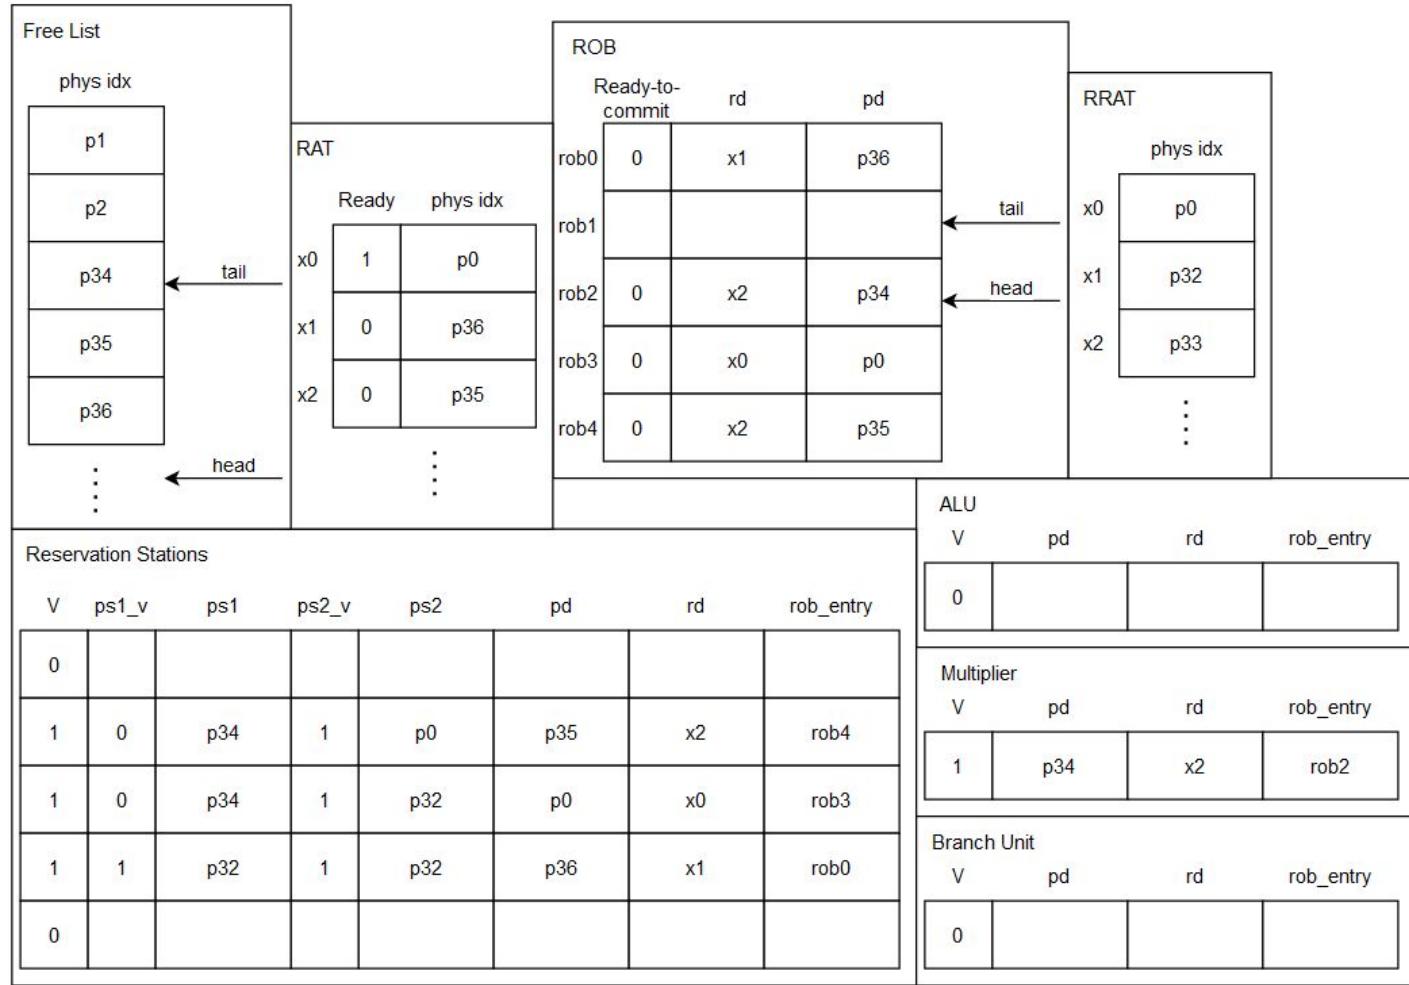


# Cycle 6:

Program:

Addi x1, x0, 4  
 Add x2, x1, x0  
 Mul x2, x2, x1  
 Bge x2, x1, label  
 Add x2, x2, x0  
 Add x1, x1, x1

Mul: 3 cycles  
 ALU: 1 cycles  
 BR: 1 cycles

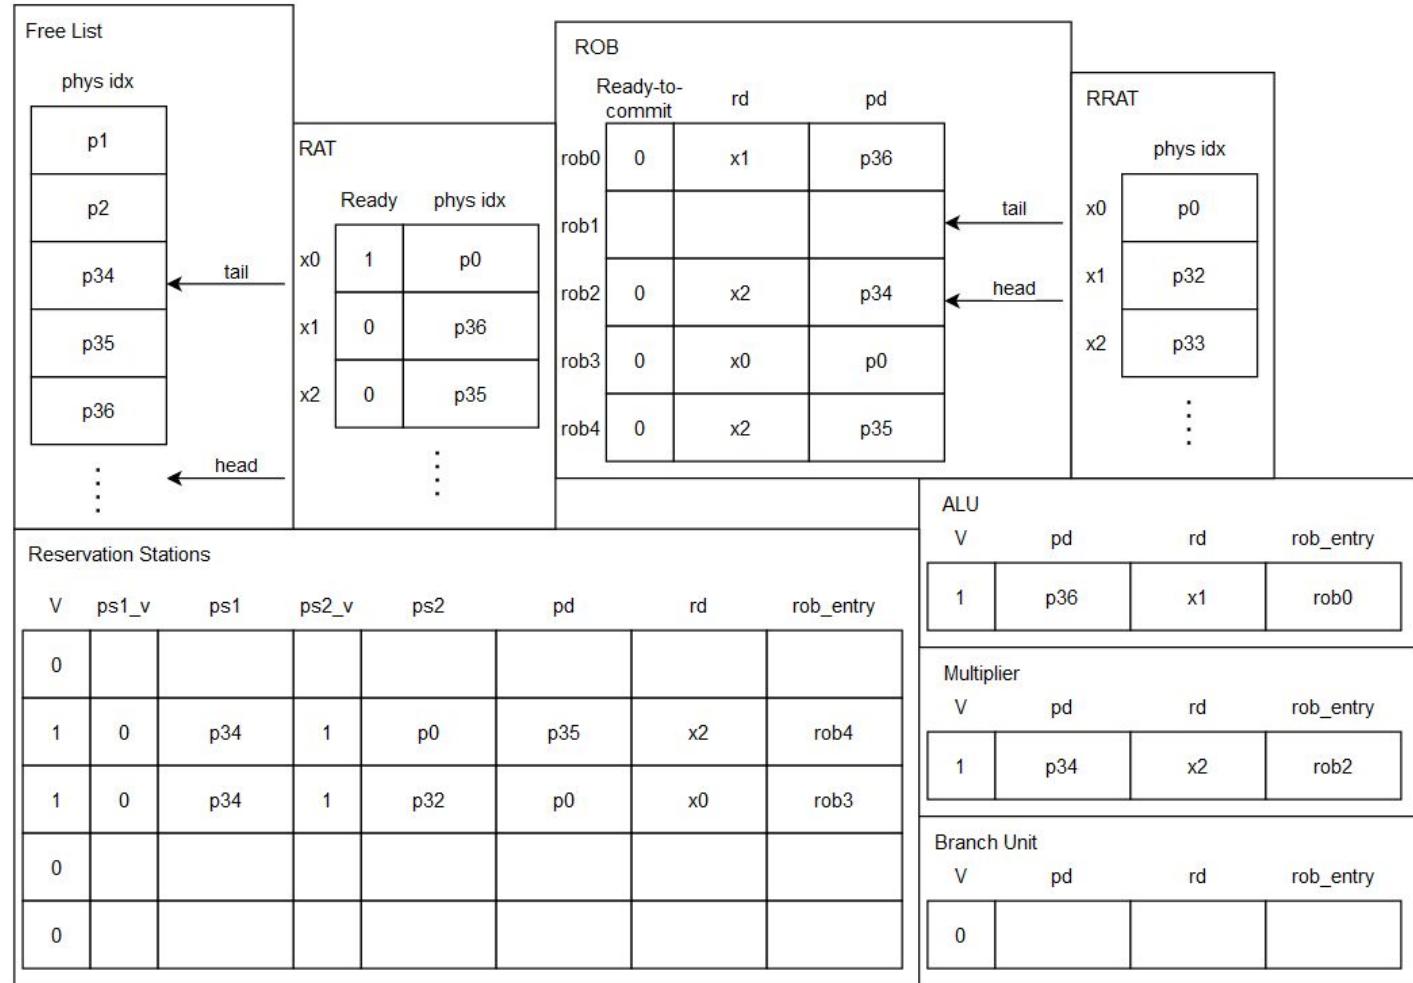


# Cycle 7:

Program:

Addi x1, x0, 4  
 Add x2, x1, x0  
 Mul x2, x2, x1  
 Bge x2, x1, label  
 Add x2, x2, x0  
 Add x1, x1, x1

Mul: 3 cycles  
 ALU: 1 cycles  
 BR: 1 cycles

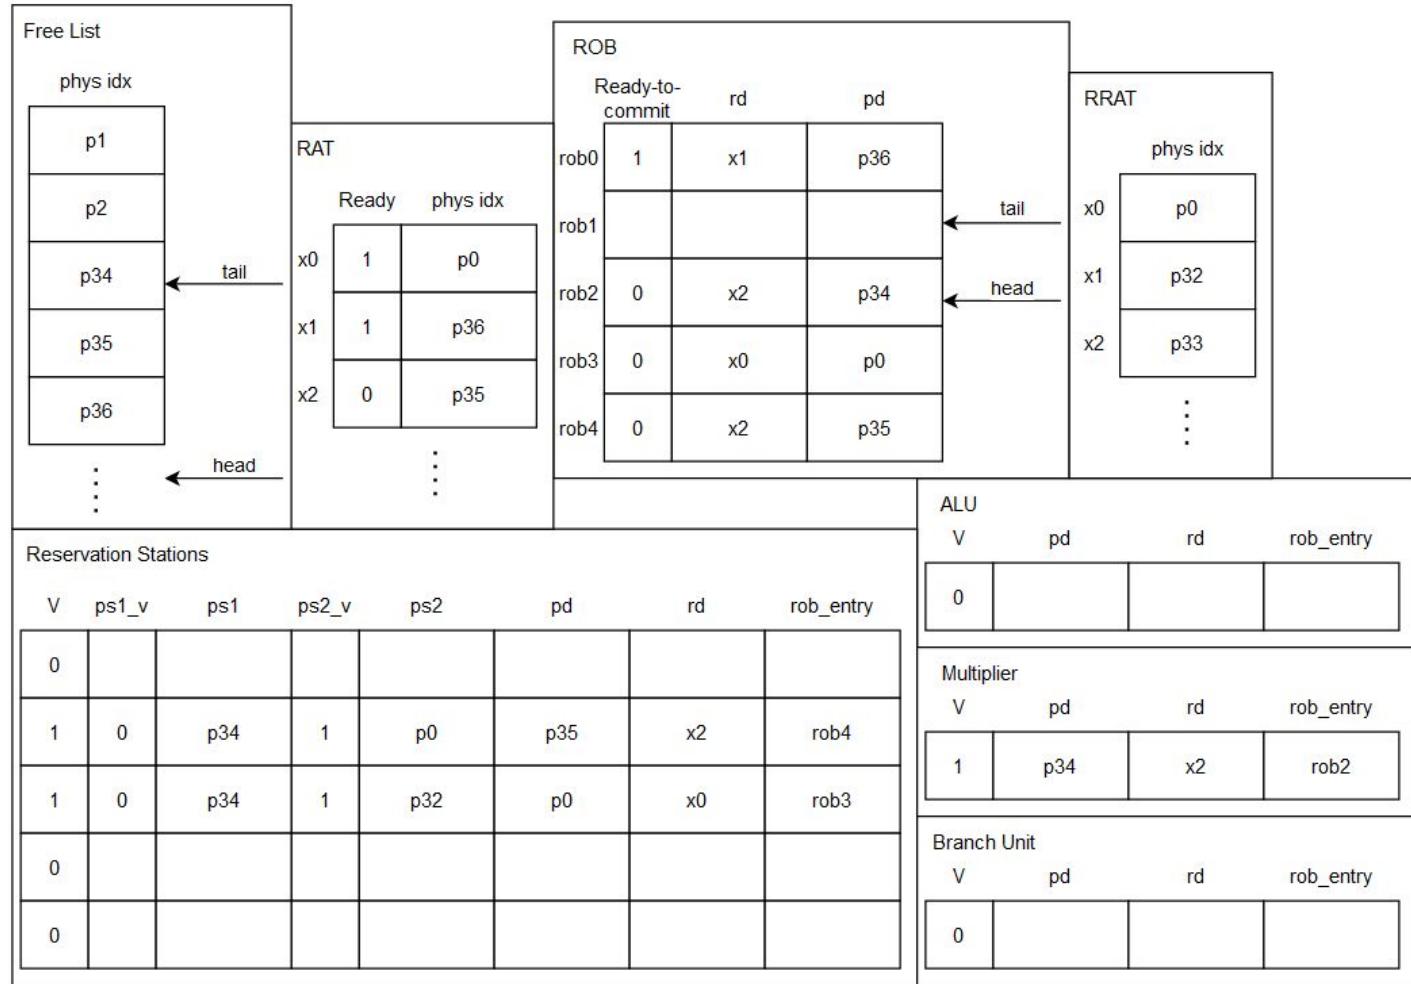


# Cycle 8:

Program:

```
Addi x1, x0, 4
Add x2, x1, x0
Mul x2, x2, x1
Bge x2, x1, label
Add x2, x2, x0
Add x1, x1, x1
```

Mul: 3 cycles  
 ALU: 1 cycles  
 BR: 1 cycles

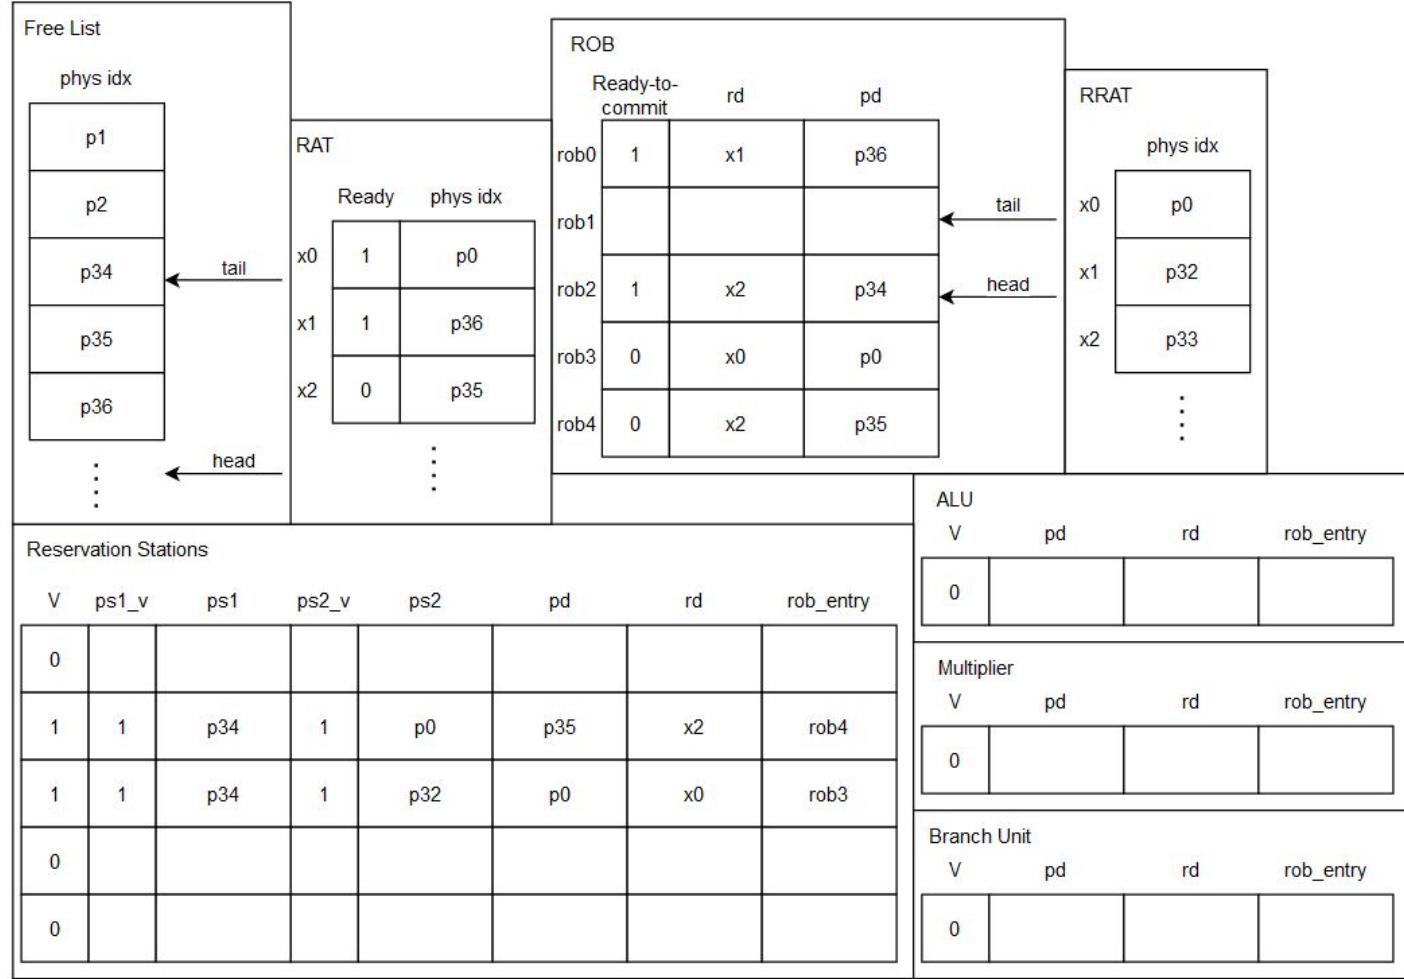


# Cycle 9:

Program:

Addi x1, x0, 4  
 Add x2, x1, x0  
 Mul x2, x2, x1  
 Bge x2, x1, label  
 Add x2, x2, x0  
 Add x1, x1, x1

Mul: 3 cycles  
 ALU: 1 cycles  
 BR: 1 cycles



# Cycle 10:

Program:

Addi x1, x0, 4

Add x2, x1, x0

Mul x2, x2, x1

Bge x2, x1, label

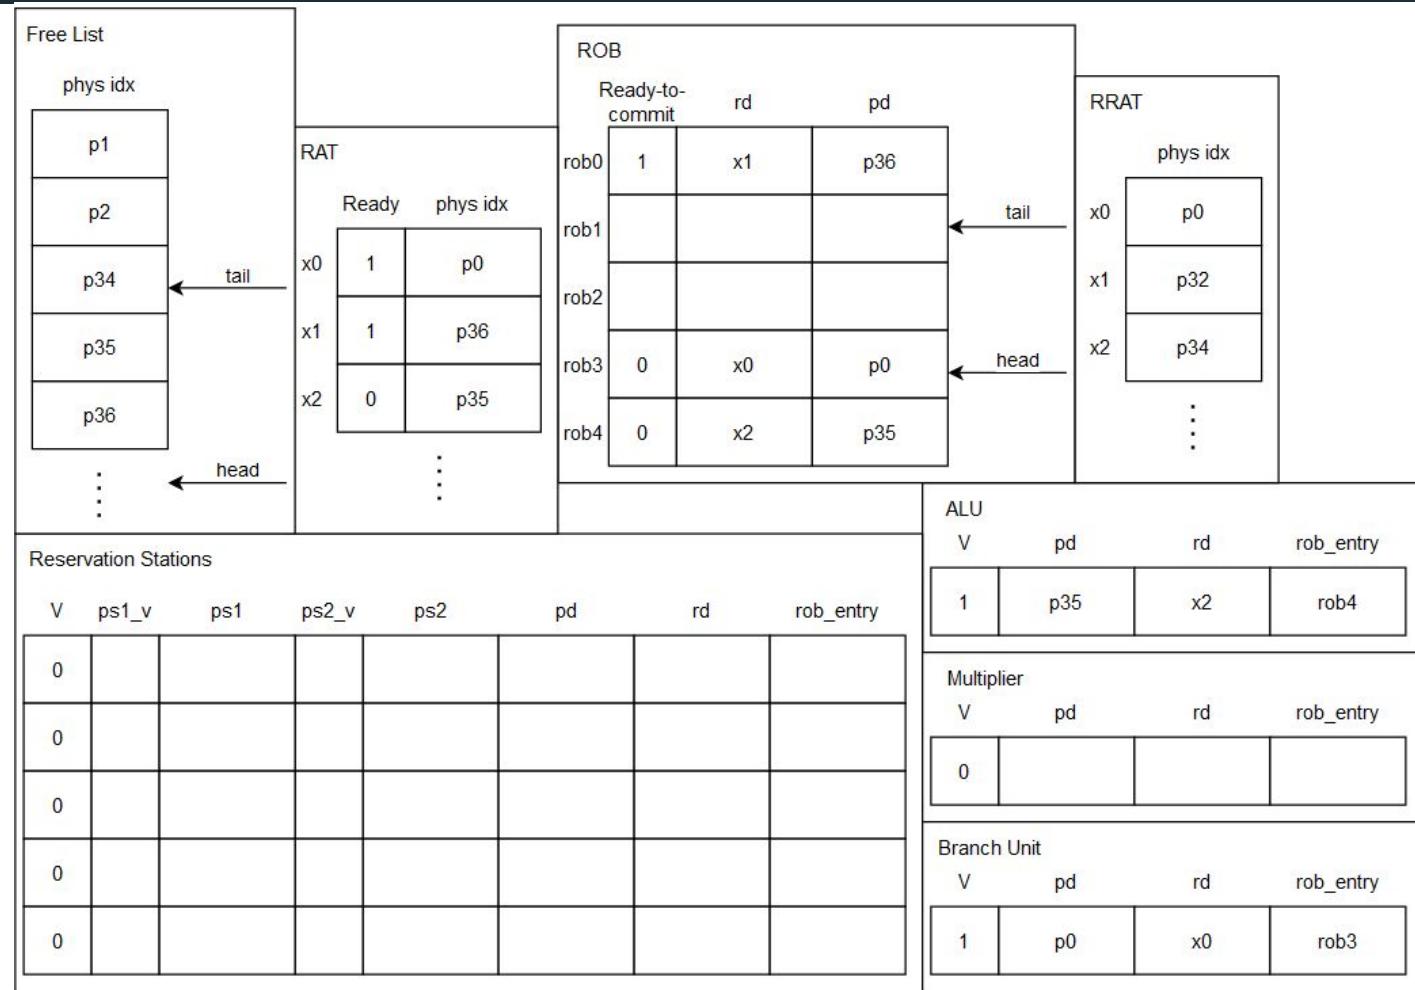
Add x2, x2, x0

Add x1, x1, x1

Mul: 3 cycles

ALU: 1 cycles

BR: 1 cycles

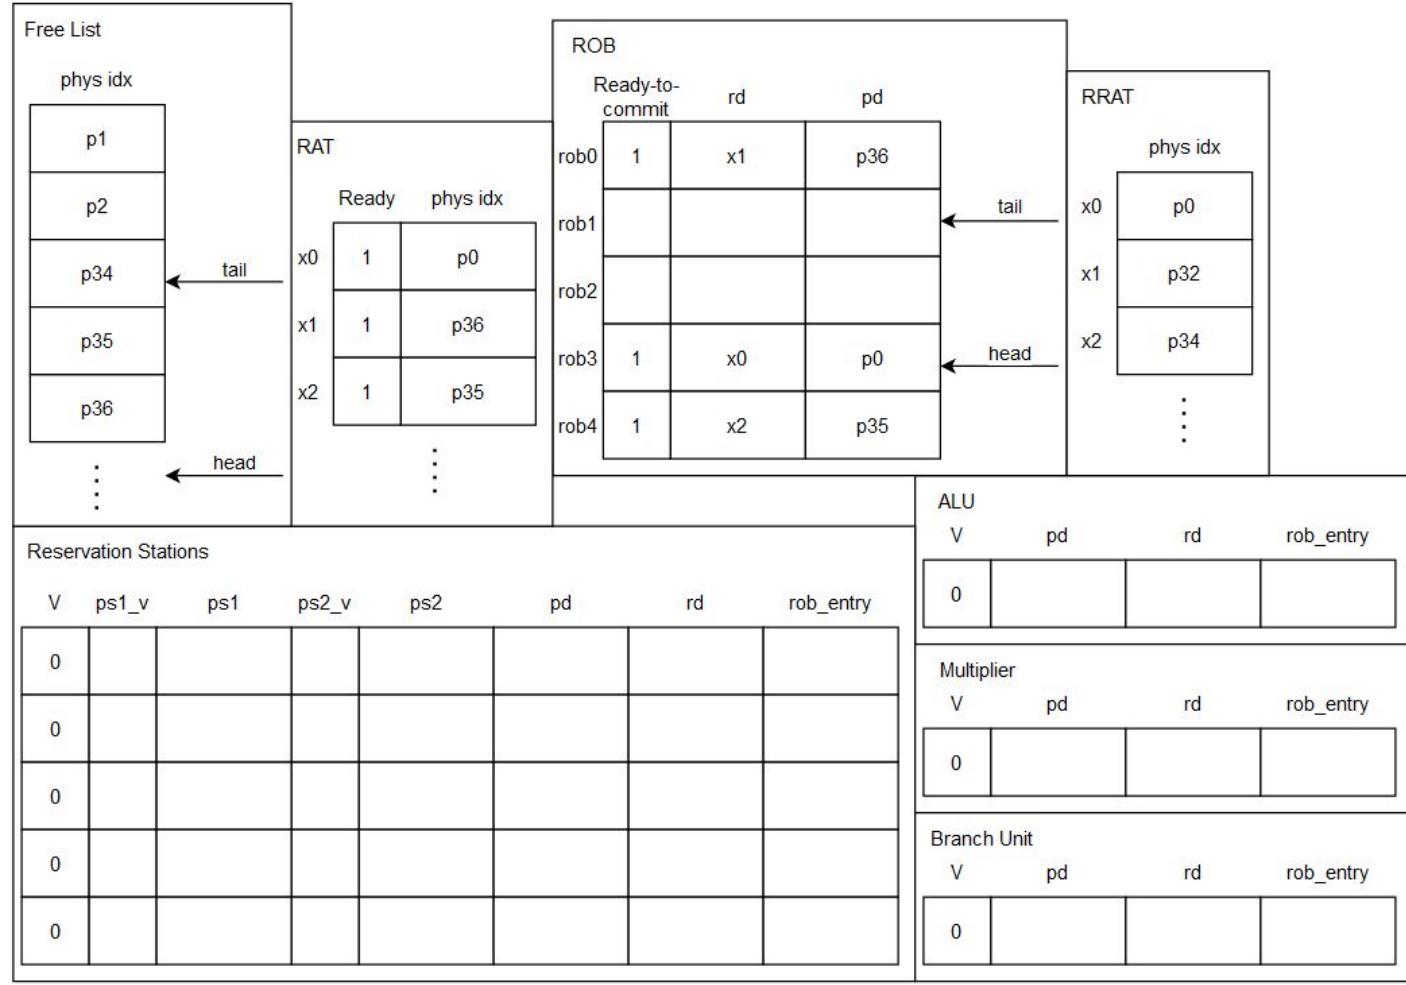


# Cycle 11:

Program:

```
Addi x1, x0, 4
Add x2, x1, x0
Mul x2, x2, x1
Bge x2, x1, label
Add x2, x2, x0
Add x1, x1, x1
```

Mul: 3 cycles  
 ALU: 1 cycles  
 BR: 1 cycles



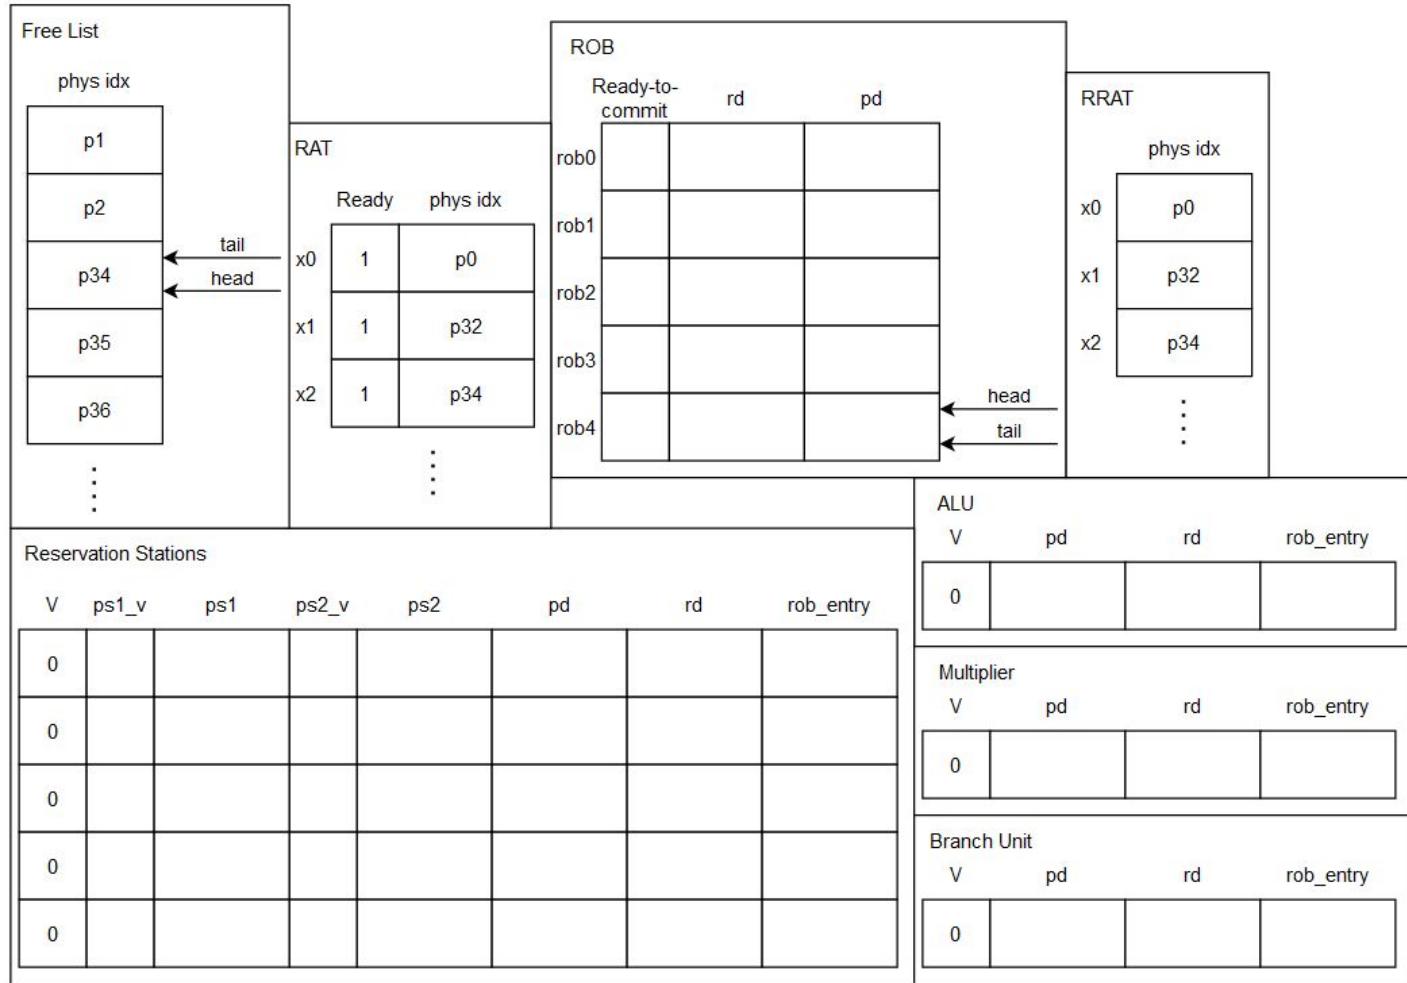
# Cycle 12:

Program:

Addi x1, x0, 4  
 Add x2, x1, x0  
 Mul x2, x2, x1  
 Bge x2, x1, label

Add x2, x2, x0  
 Add x1, x1, x1

Mul: 3 cycles  
 ALU: 1 cycles  
 BR: 1 cycles



# Memory Instructions in ERR

You cannot mix memory instructions with normal instructions inside the reservation stations

You need a separate structure (load-store queue, etc) for memory instructions.



# GOOD LUCK!

You're smart, you've got this,  
and we're here to help.



# Alternate Verification Tooling

# What is Verilator?

- Like VCS, converts your hardware to C++. But it's open source!
  - Can run on your local machines as a result, and is cross-platform
- Drawbacks: there's a lot.
  - Cycle-level simulation
  - *Only supports dual-state simulation*
  - Also pickier about the quality of your HDL (stricter linting)
  - Incomplete SV support compared to VCS
  - Harder to support threaded models
  - Can't access internal signals by default
  - Traces are way bigger compared to VCS (15GB for all of Coremark with struct extraction)
- Why bother using it at all?
  - ***Speed - “all of CoreMark within seconds” type speed.***

# Verilator Testbenches

- Overall simulation logic is no longer managed by an SV initial block
- Need to write C++ logic to interface with the DUT, manage CLK, etc.
  - Can still write performance monitoring in SV
- State management in C++ -> powerful software models in a familiar language
  - Lack of bitslices can sometimes make memory access challenging
- Many other features, so definitely look at the documentation on their website!